

# Google Native Client

## Analysis Of A Secure Browser Plugin Sandbox

Chris Rohlf  
Leaf SR

# Me

- Chris Rohlf
- Founder: Leaf [SR](#)
- BlackHat Review Board & Past Speaker 2009/2011
- @chrisrohlf
- Chris.Rohlf@gmail.com
- leafsr.com



# Google Native Client

browsers are the new platform for applications

# Google Native Client

browsers are the new platform for applications  
sandboxes are the future of application security



# Google Native Client

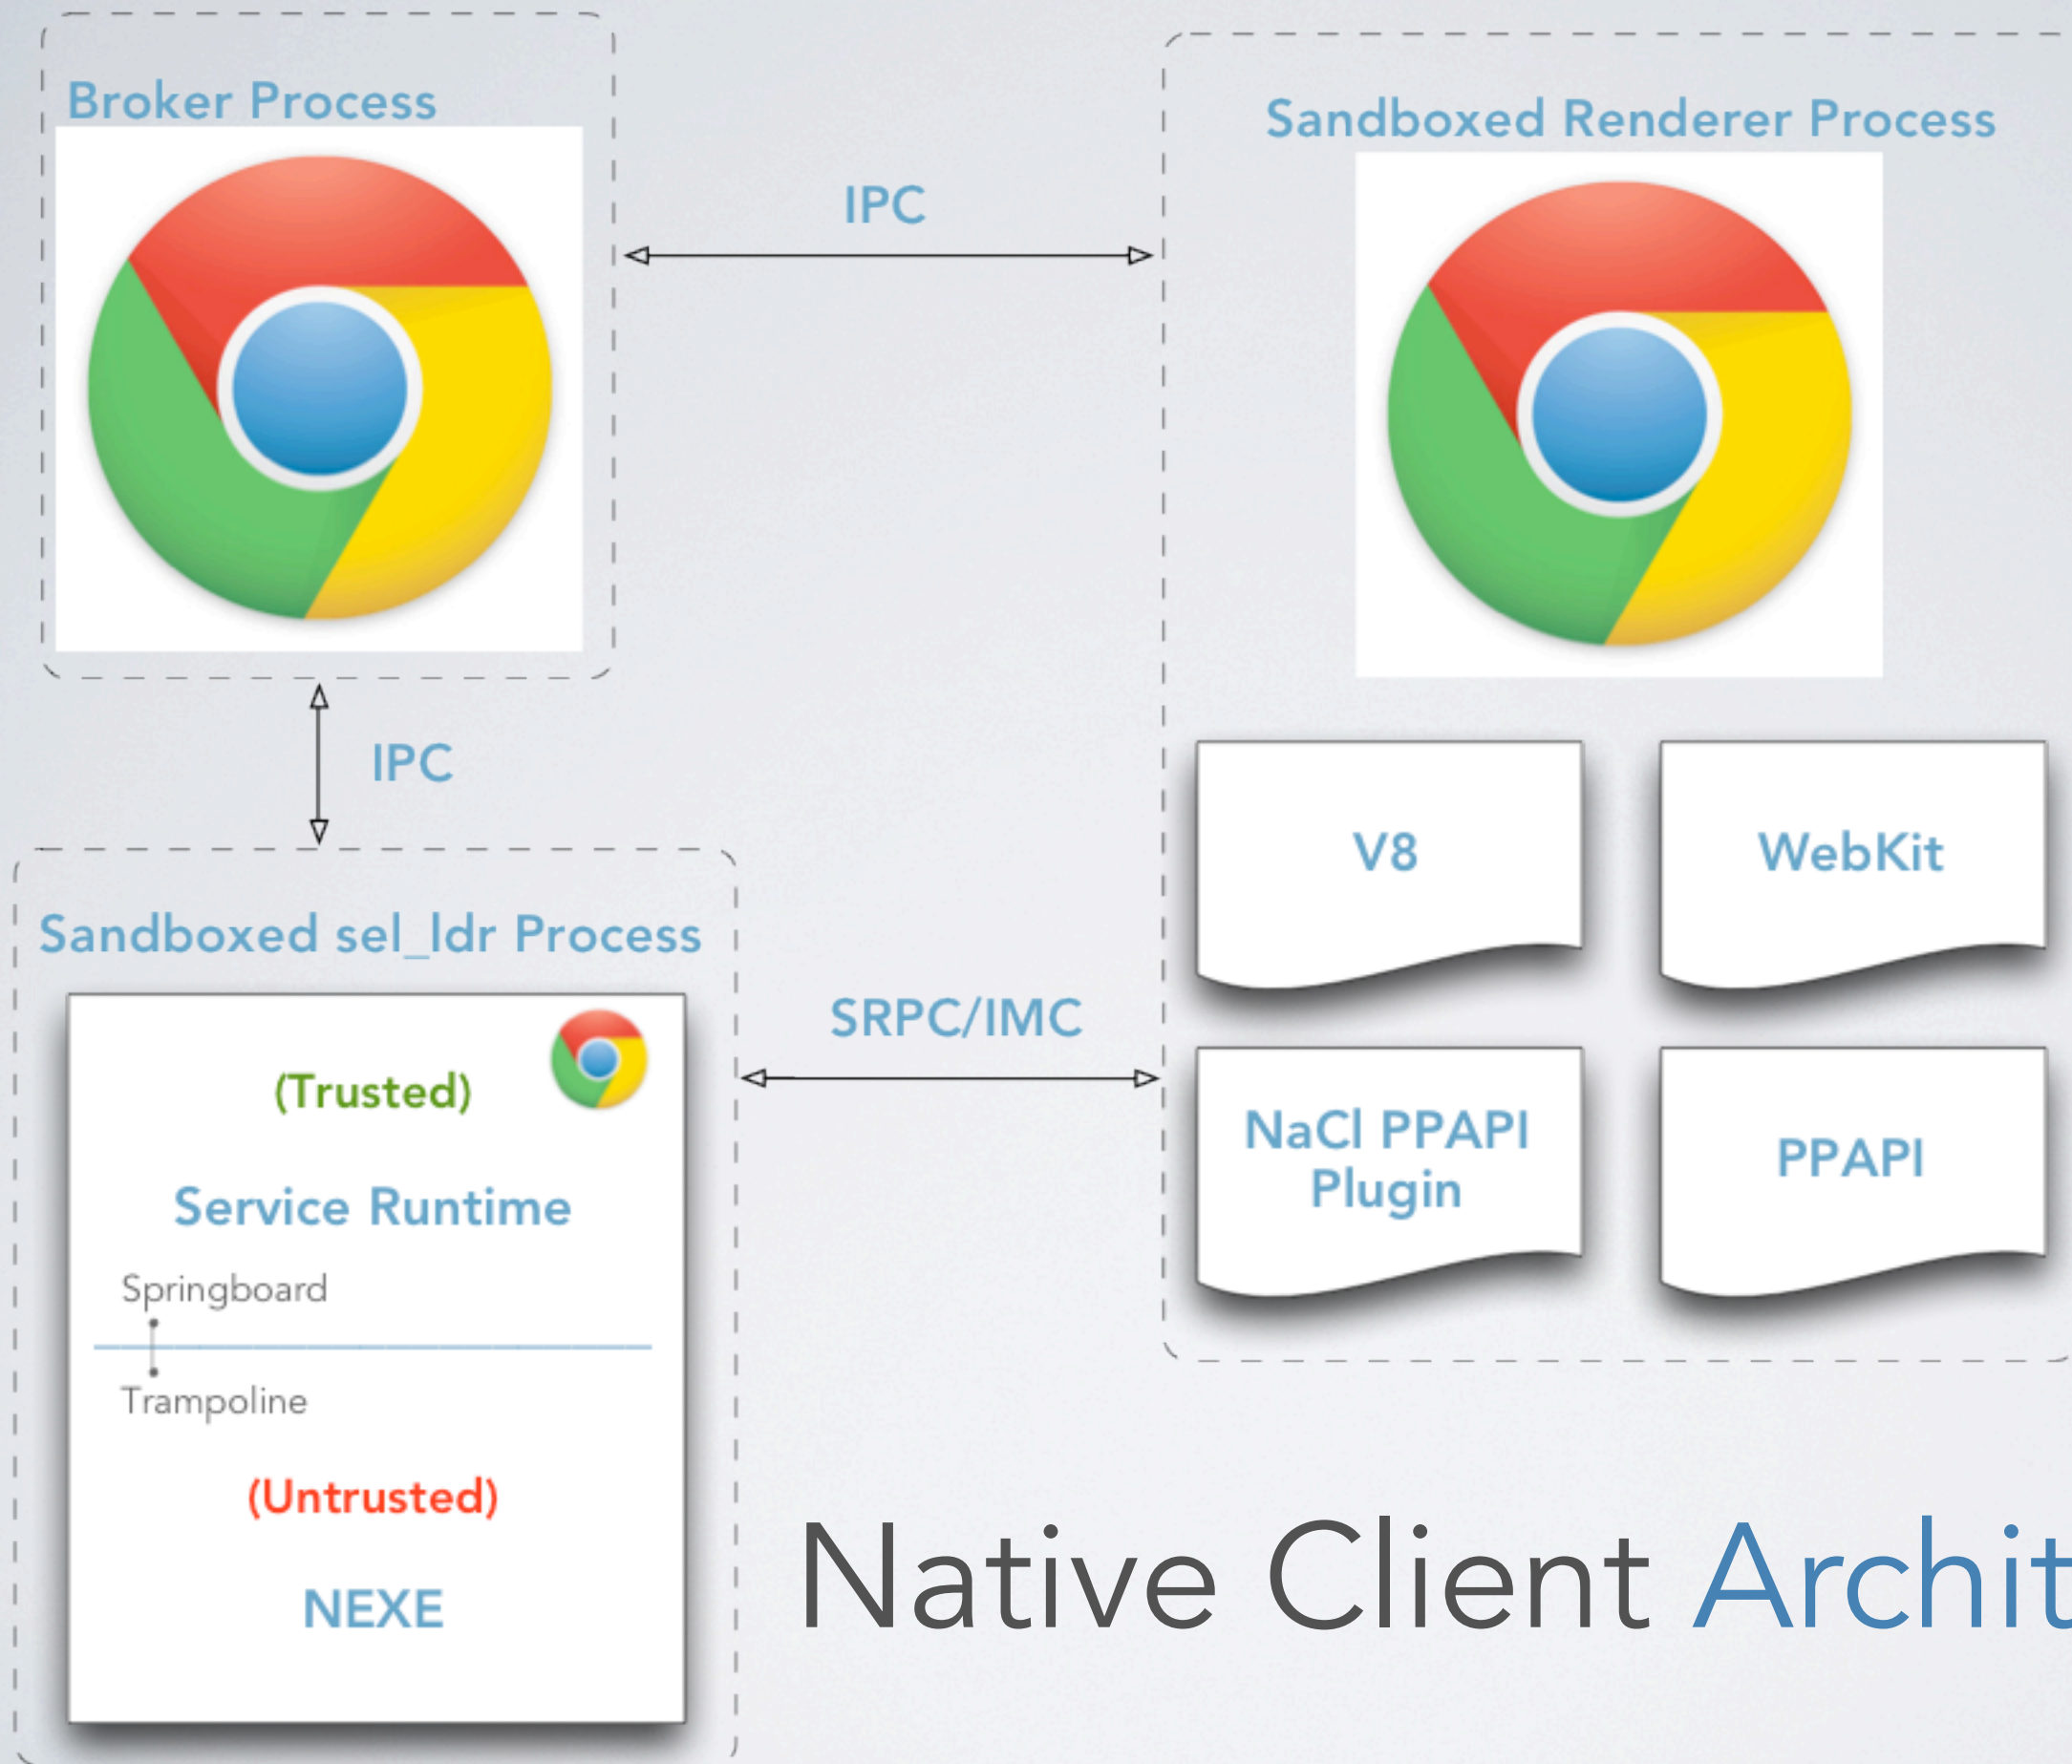
browsers are the new platform for applications  
sandboxes are the future of application security  
breaking easy targets is boring!



# Google Native Client

- A Chrome plugin that allows the execution of native untrusted code in your browser (Win32, OS X, Linux)
- Chrome 14 shipped with NaCl by default
- Large and complex architecture
  - Modified Compiler Toolchain
  - Secure ELF loader
  - Disassembler and Code Validator
  - Service Runtime
  - Inner and Outer Sandbox
  - SRPC (Simple Remote Procedure Call)
  - IMC (Inter-Module Communication)
  - PPAPI (Pepper)



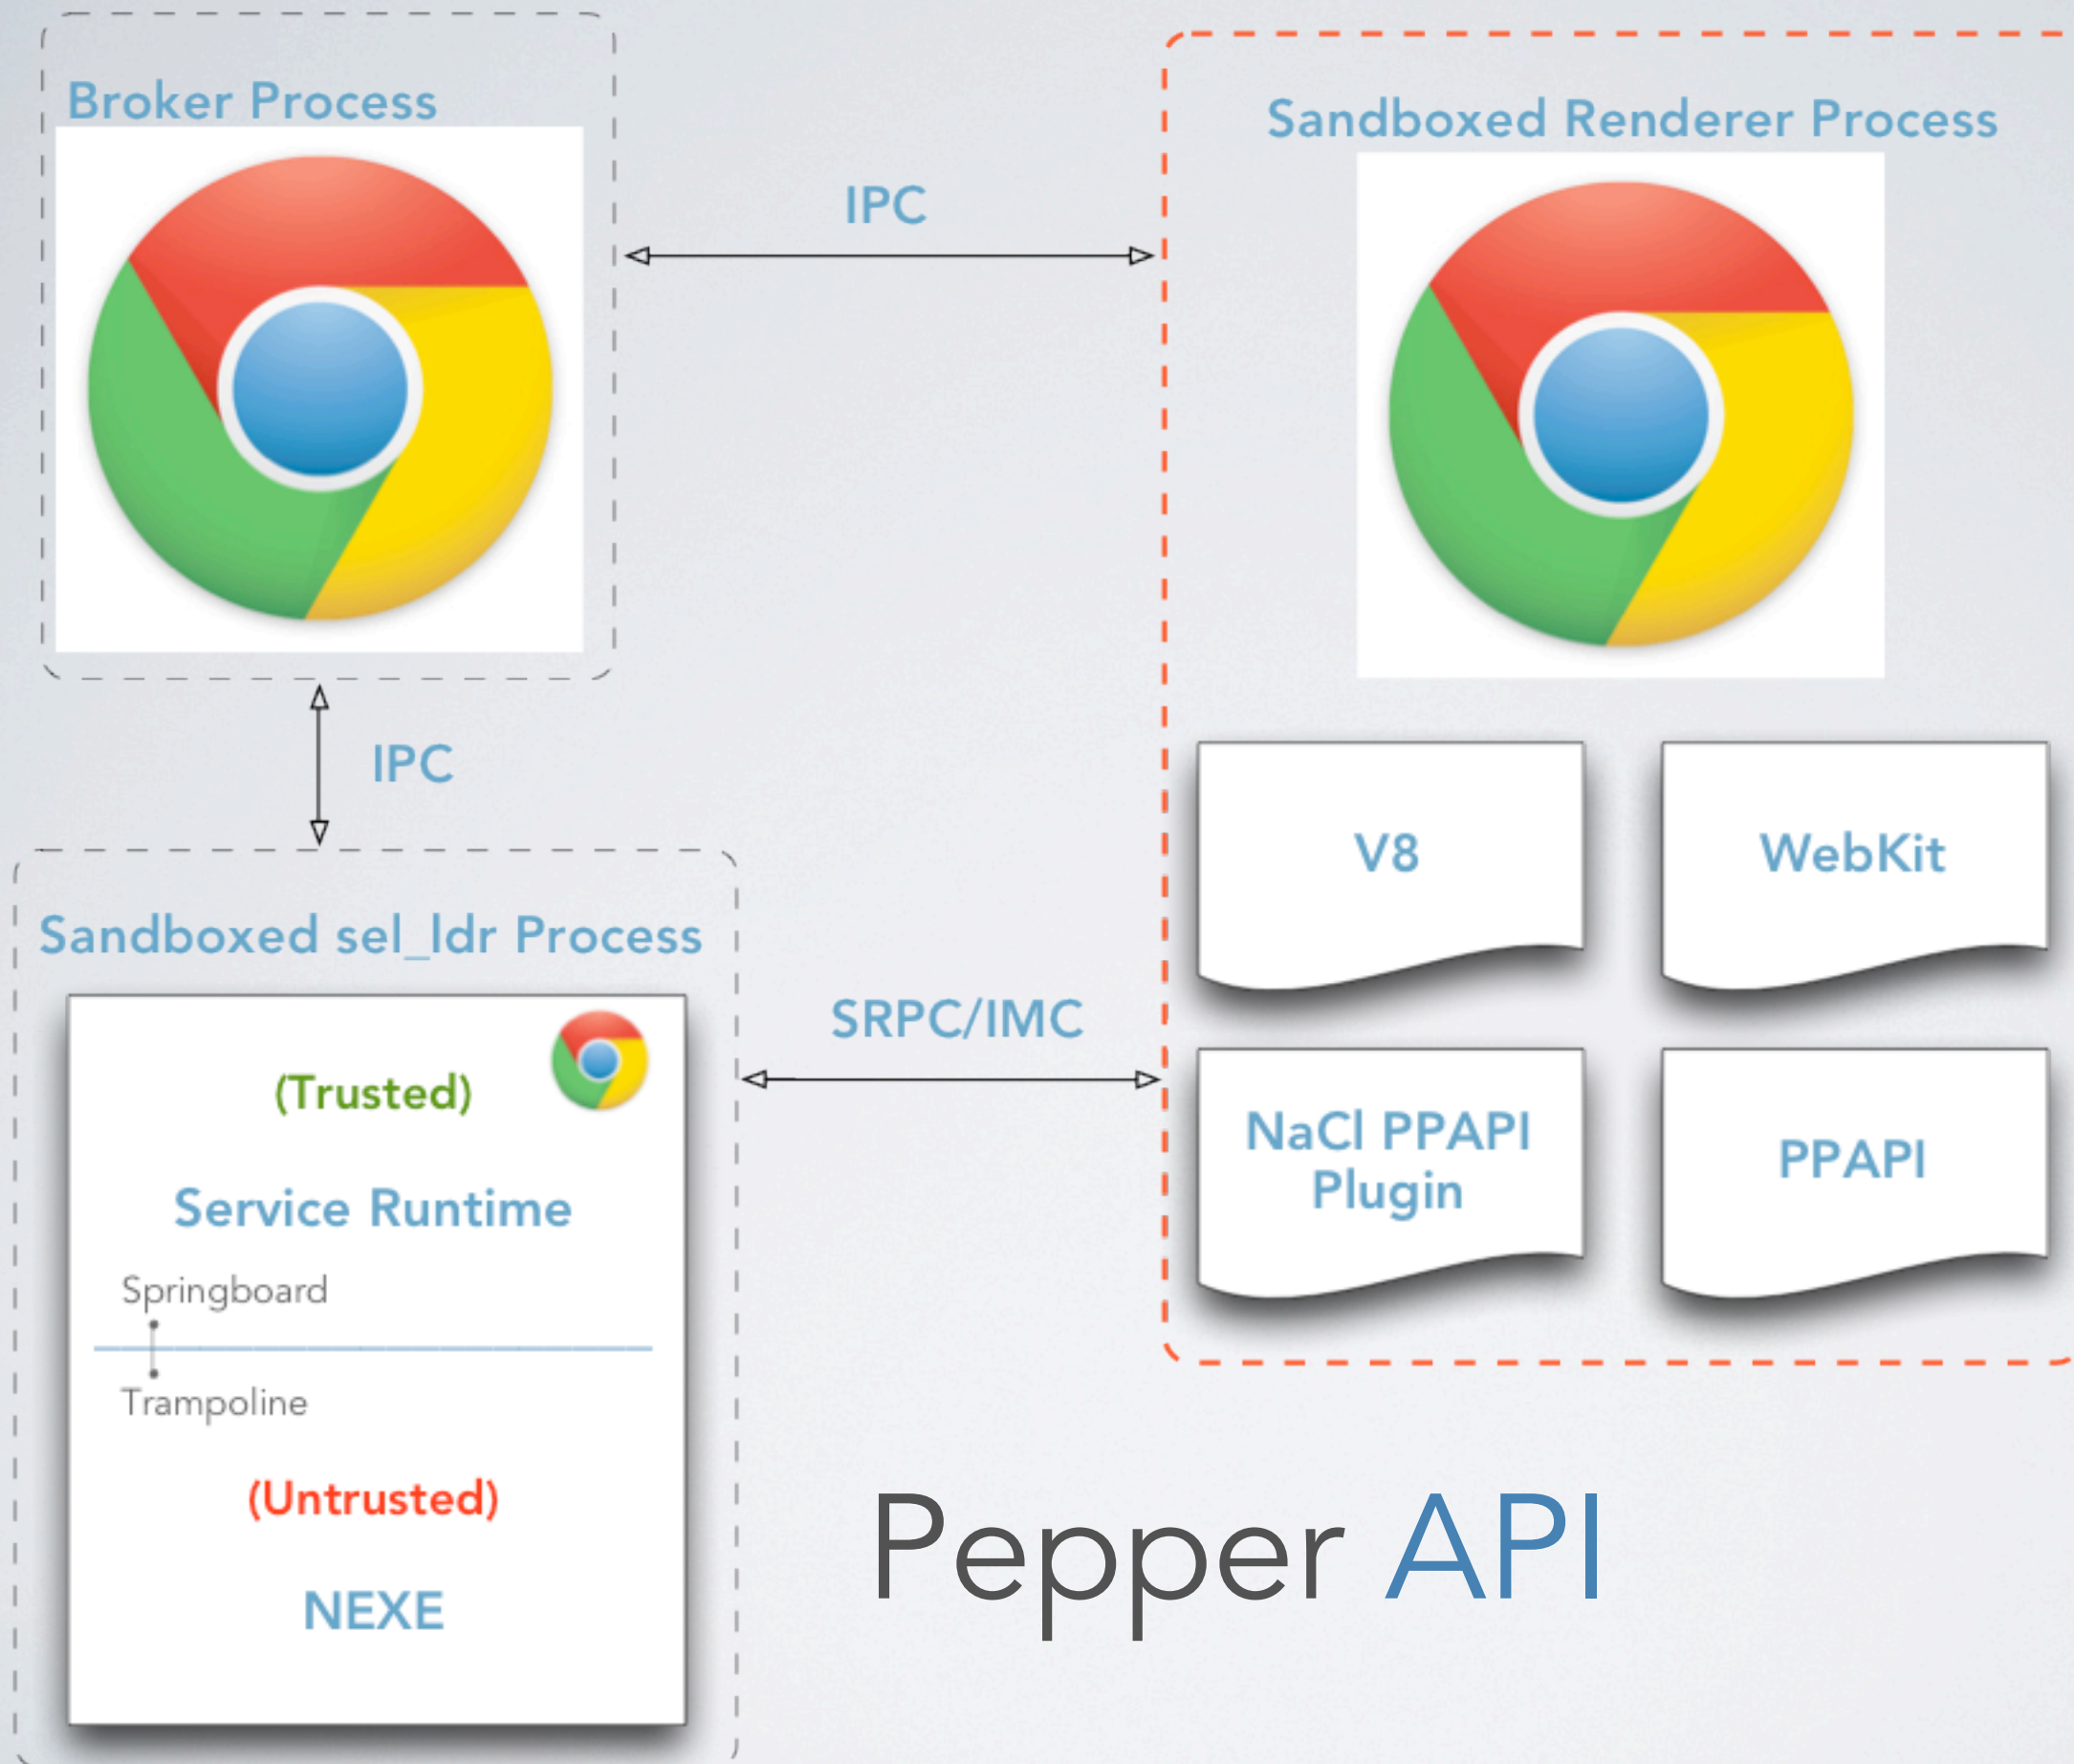


# Native Client Architecture

# Levels Of Trust

- Outer Sandbox
  - Chrome Broker - Trusted
  - Chrome Renderer - Untrusted
  - NaCl Service Runtime - Untrusted
  - NEXE Module - Untrusted
- Inner Sandbox
  - Chrome Broker - Trusted
  - Chrome Renderer - Trusted
  - NaCl Service Runtime - Trusted
  - NEXE Module - Untrusted





# Pepper API

# Pepper API

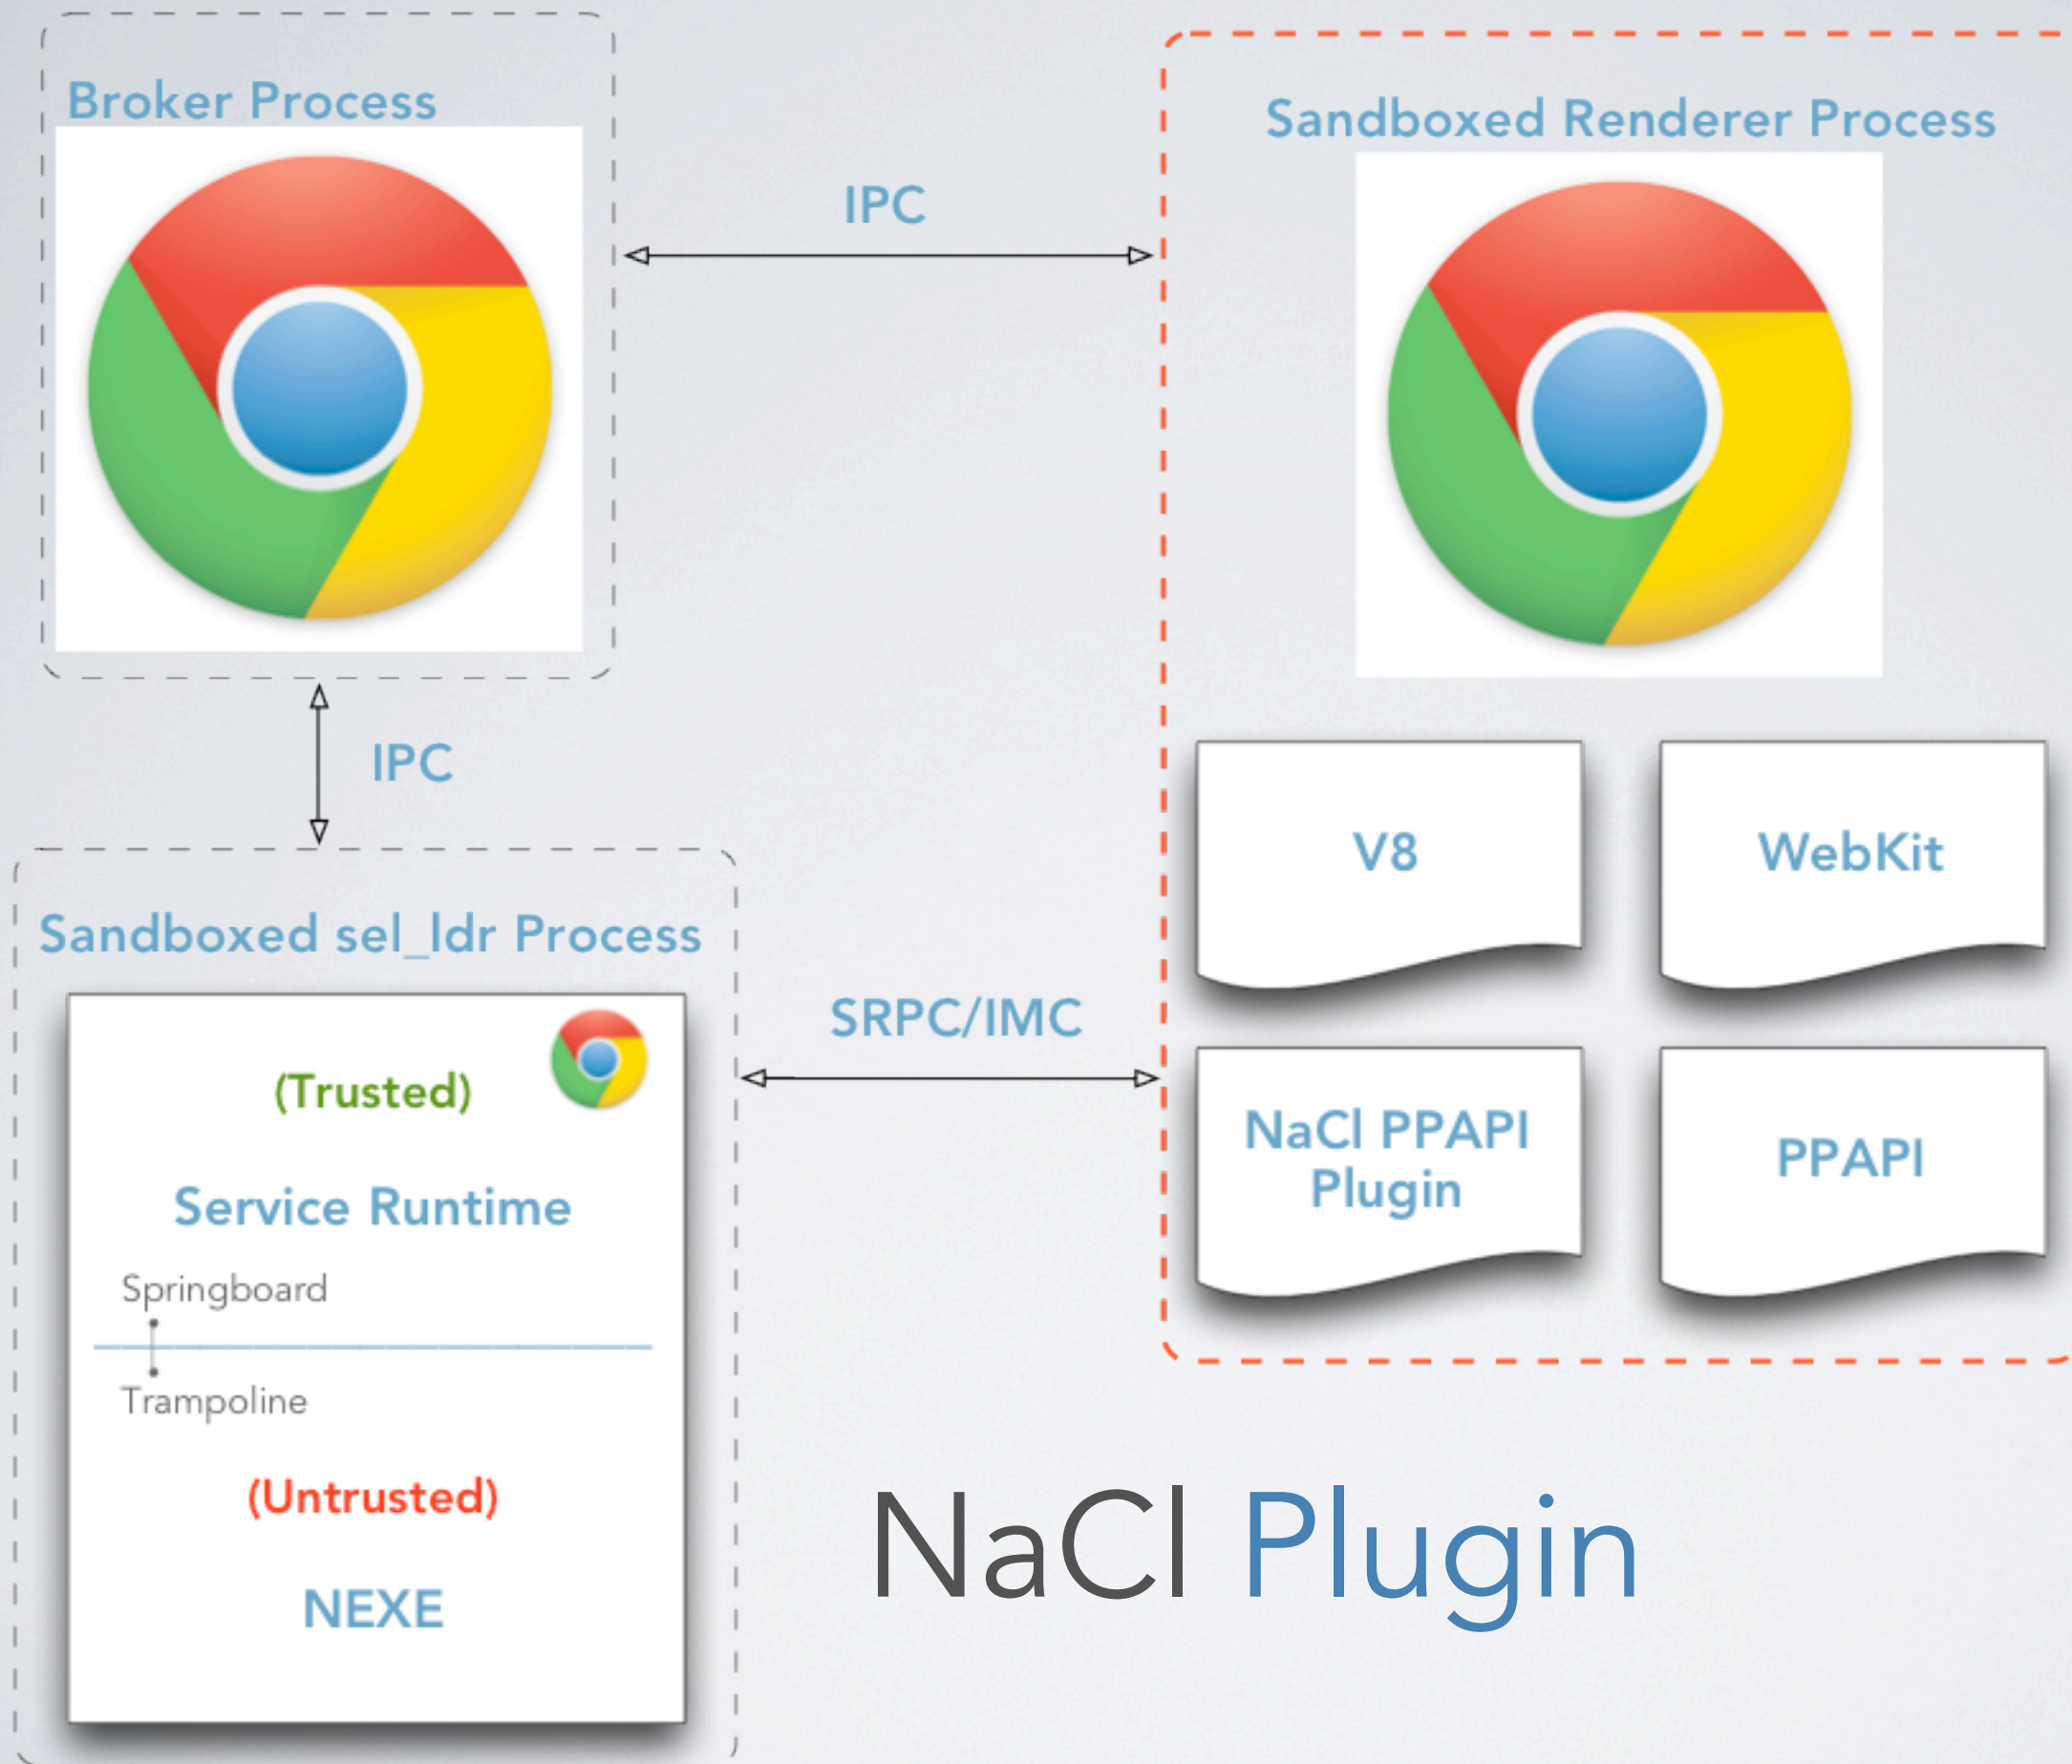
- PPAPI (Pepper Plugin Application Programming Interface)
  - Replaces NPAPI (Netscape Plugin API)
- New APIs for Audio, 3D, Input devices and more
- Interfaces for privileged actions (FileIO...)
- Not scriptable via JavaScript
- Out of process plugins are supported by the standard



# PPAPI Plugins

- Trusted PPAPI plugins run in a sandboxed Chrome renderer process or a separate sandboxed plugin process
  - Native Client is a trusted plugin
  - Adobe Pepper Flash is a trusted out of process plugin
- Untrusted PPAPI plugins in Chrome run as a NEXE module
  - NEXE's communicate with PPAPI layer using a proxy





# NaCl Plugin



# NaCl Plugin

- The NaCl plugin itself is a trusted PPAPI plugin
- Lives inside the Chrome renderer process as a DLL
- Invoked by HTML embed tag

```
<embed name="NaCl_module"  
      id="hello_world"  
      width=200 height=200  
      src="hello_world.nmf"  
      type="application/x-NaCl" />
```



# NaCl Plugin

- The .nmf file is a 'NaCl Manifest File'
- JSON that specifies NEXE and .so libraries

```
{ "files": {  
  "libgcc_s.so.1": { "x86-32": { "url": "lib32/libgcc_s.so.1" } },  
  "main.nexe": { "x86-32": { "url": "hw.nexe" } },  
  "libc.so.3c8d1f2e": { "x86-32": { "url": "lib32/libc.so.3c8d1f2e" } },  
  "libpthread.so.3c8d1f2e": { "x86-32": { "url": "lib32/libpthread.so.3c8d1f2e" } } },  
  "program": { "x86-32": { "url": "lib32/runnable-ld.so" } }  
}
```

- The plugin parses this with *jsoncpp*
- NaCl downloads the NEXE using the PPAPI URLLoader and FileIO interfaces



# NaCl Plugin

- NaCl must expose part of itself to the browser DOM
- However PPAPI is not scriptable like NPAPI
- Only a few properties and functions are exposed
  - readyState, lastError, exitStatus
  - postMessage



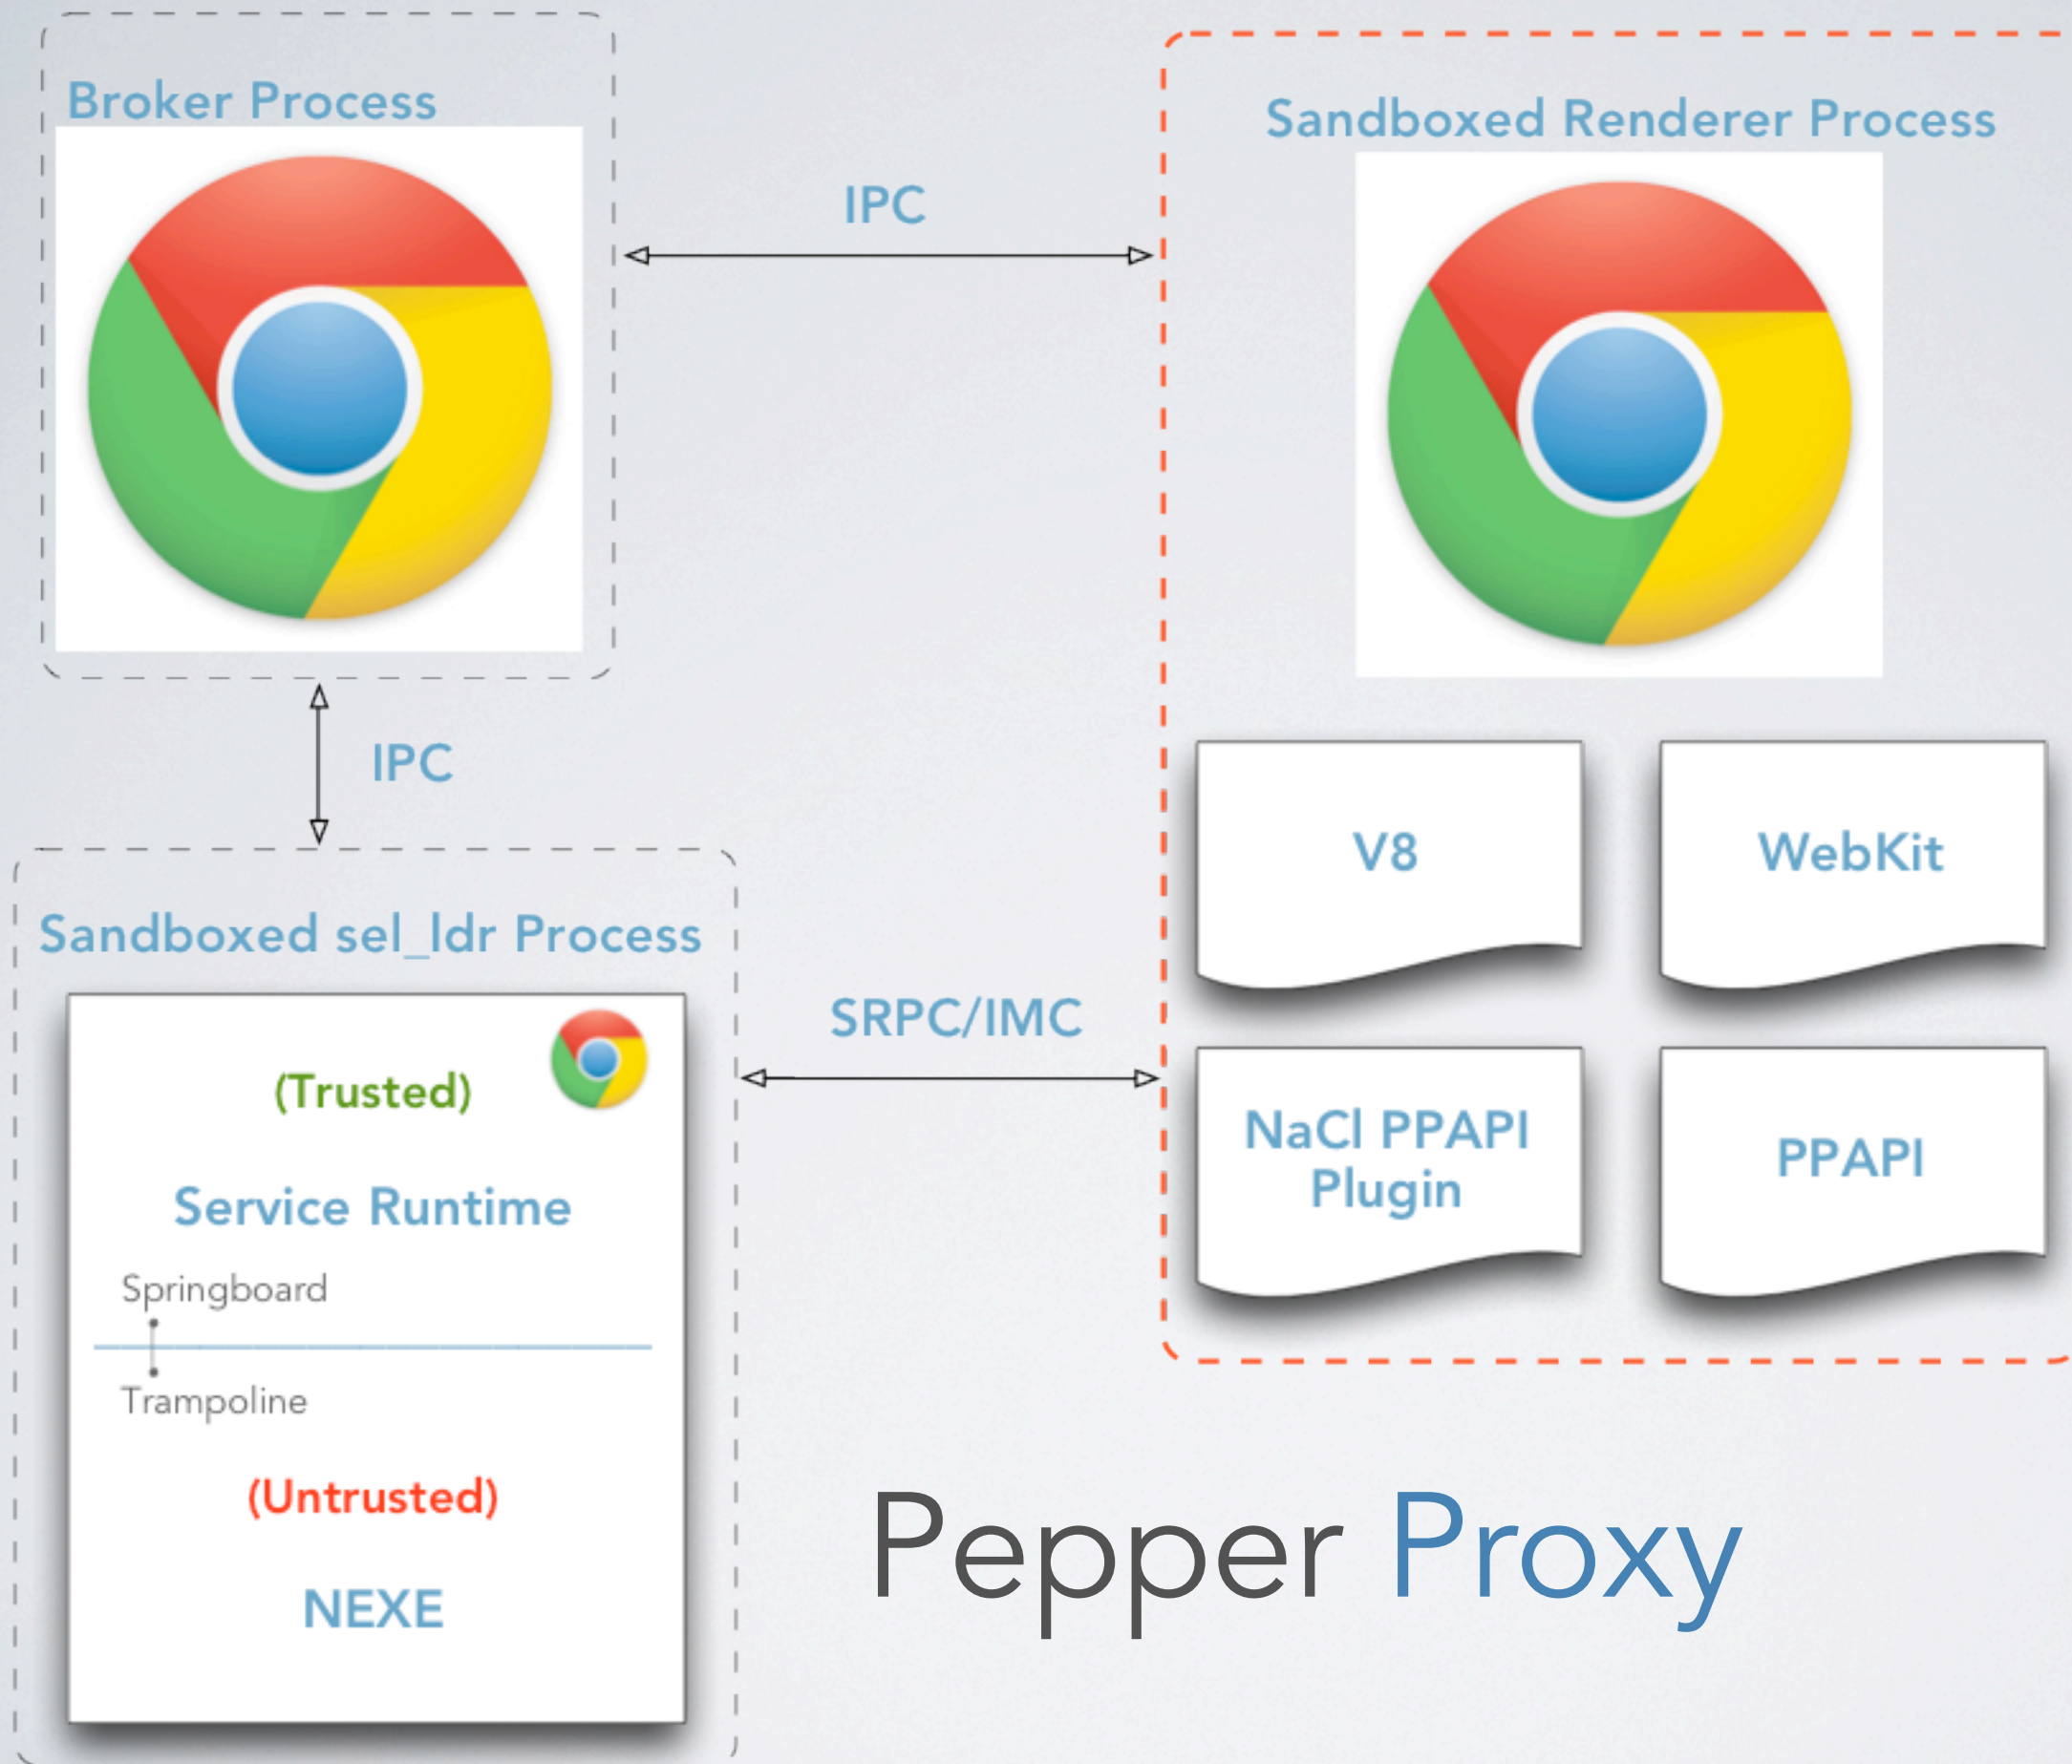
# NaCl Plugin

- PostMessage is how JavaScript talks to the NEXE module
- The NEXE must implement `pp::Instance::HandleMessage` to receive these messages
- The `PPB_Messaging::PostMessage` C++ interface is used to send messages back to JavaScript
- JavaScript gets alerted to new messages using an event listener and a callback
- These messages have no format and can be binary or ASCII



# NaCl Plugin

- NaCl asks the broker to start the service runtime process
- Once the service runtime has started they establish an SRPC connection with each other
- This is the administrative SRPC channel
- Individual SRPC channels will be created for each untrusted NEXE module that is launched



# Pepper Proxy



# Pepper Proxy

- NEXE modules need to talk to the browser in order to do anything useful
- PPAPI provides plugins the ability to access privileged browser resources
- The NaCl plugin implements the pepper proxy
- This is the bridge between PPAPI and the untrusted NEXE



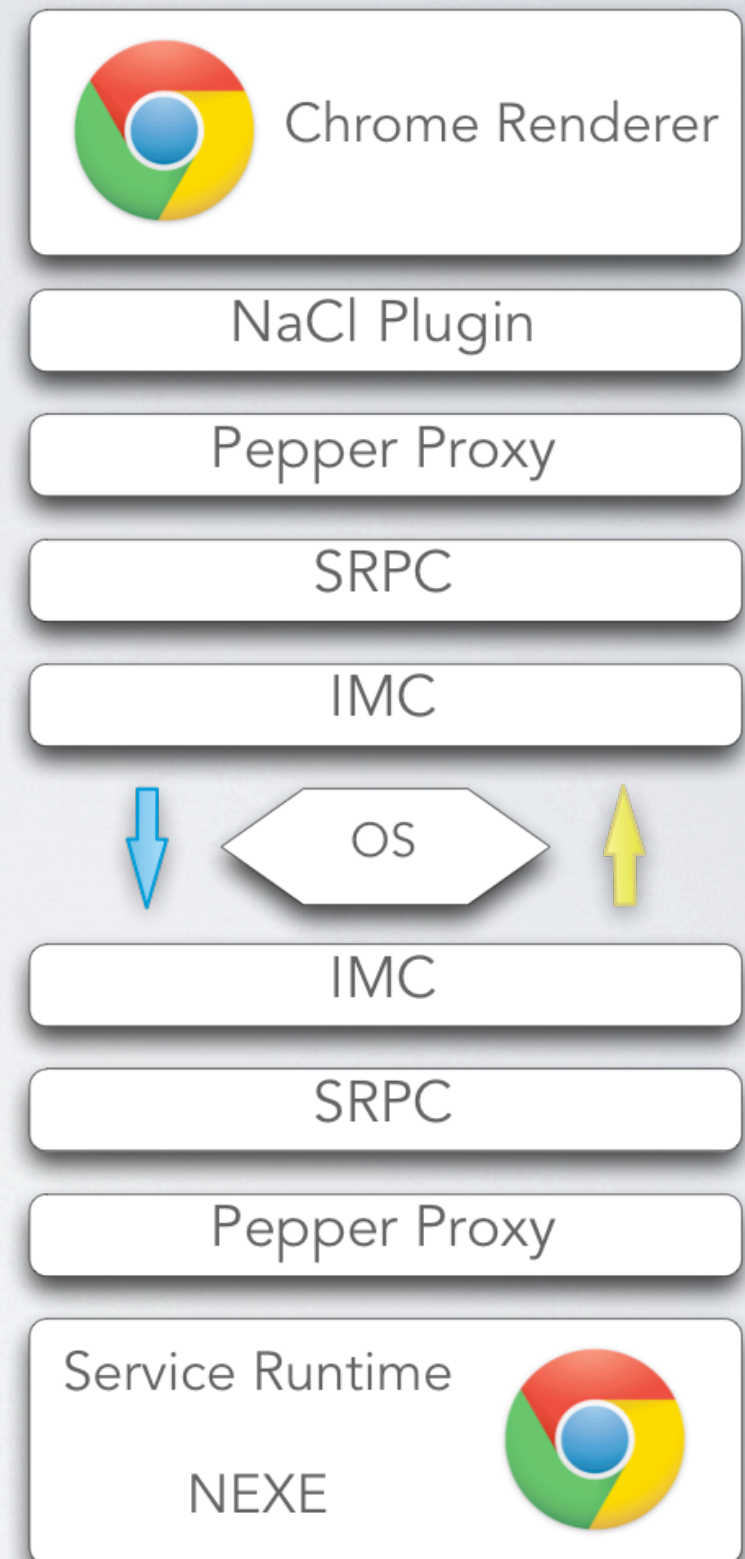
# Pepper Proxy

- NaCl uses IDL files to describe these interfaces
  - Pepper proxy code is auto generated using them
- Both sides can act as server and/or client
  - Shared code between trusted/untrusted sides
- A simple rule to help differentiate the two
  - Interfaces prefixed with PPP are on the untrusted side
  - Interfaces prefixed with PPB are on the trusted side



# Pepper Proxy

- The protocol stack allows untrusted NEXE modules to invoke trusted PPAPI interfaces
- Serialized PPAPI arguments, over SRPC, over IMC
- Most of this is binary data packaged up as PP\_Var or basic data types





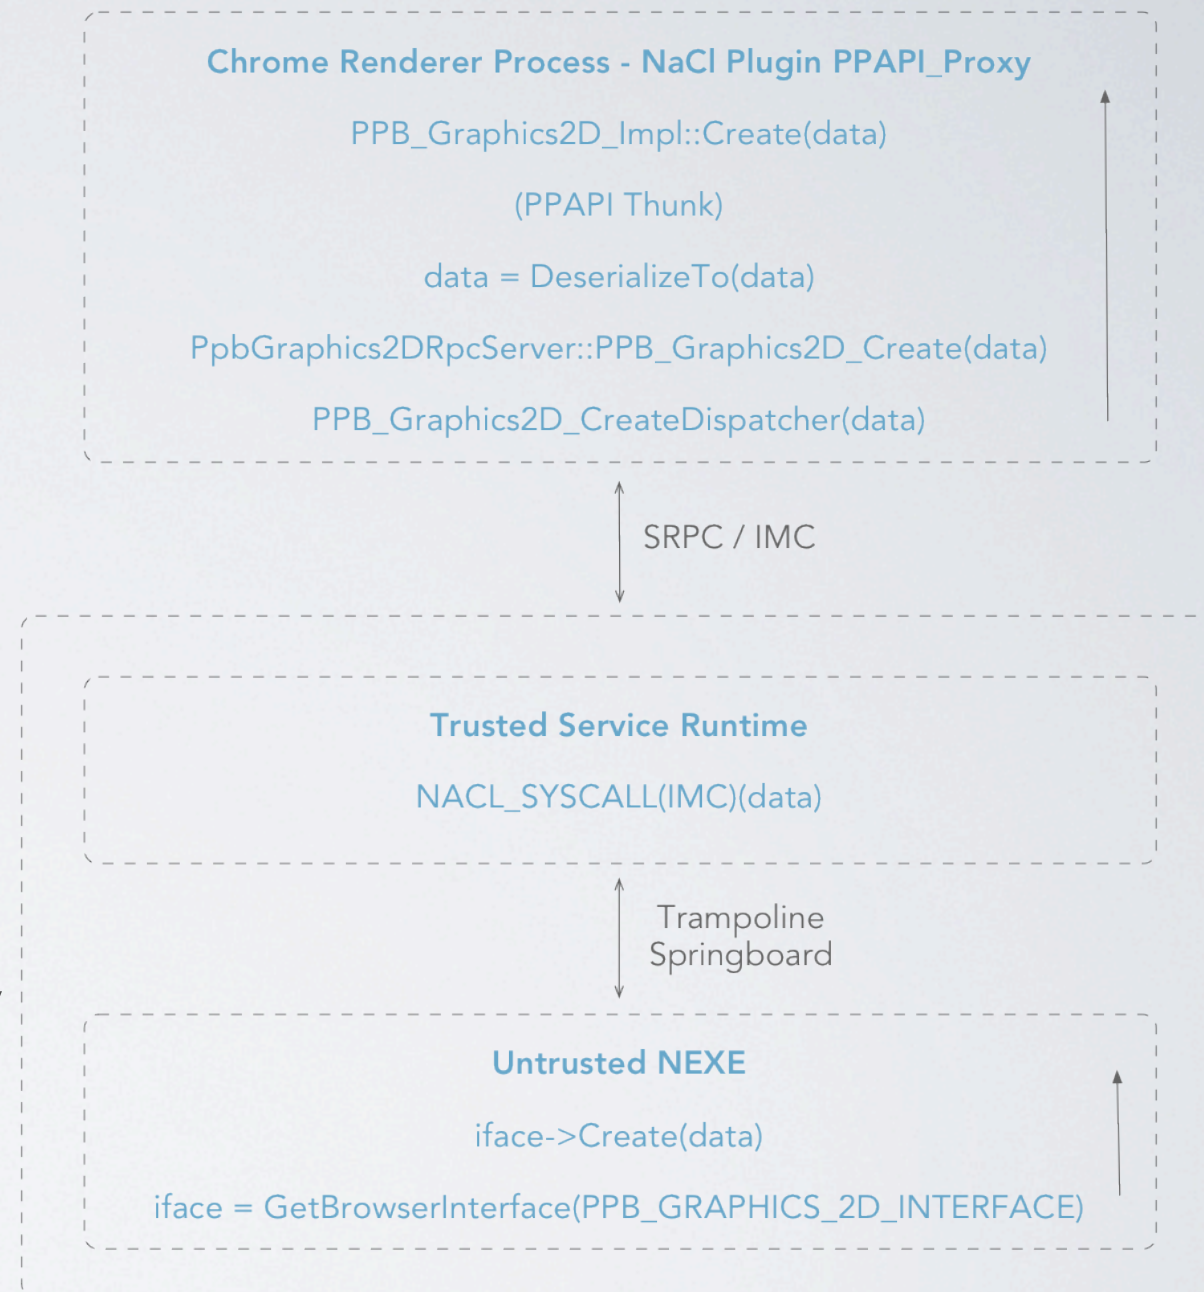
# Pepper Proxy

- Data sent over the proxy is serialized
  - PP\_VarType - enum specifying type of PP\_Var
  - PP\_VarValue - union holding simple data types
  - PP\_Var - structure that holds PP\_VarType and PP\_VarValue



# Pepper Proxy

1. Creates a PP\_Var and other args
2. Calls a remote pepper interface
3. SRPC message gets created
4. SRPC message sent over IMC
5. Received by remote pepper proxy
6. Data deserialized
7. Passed off to PPAPI





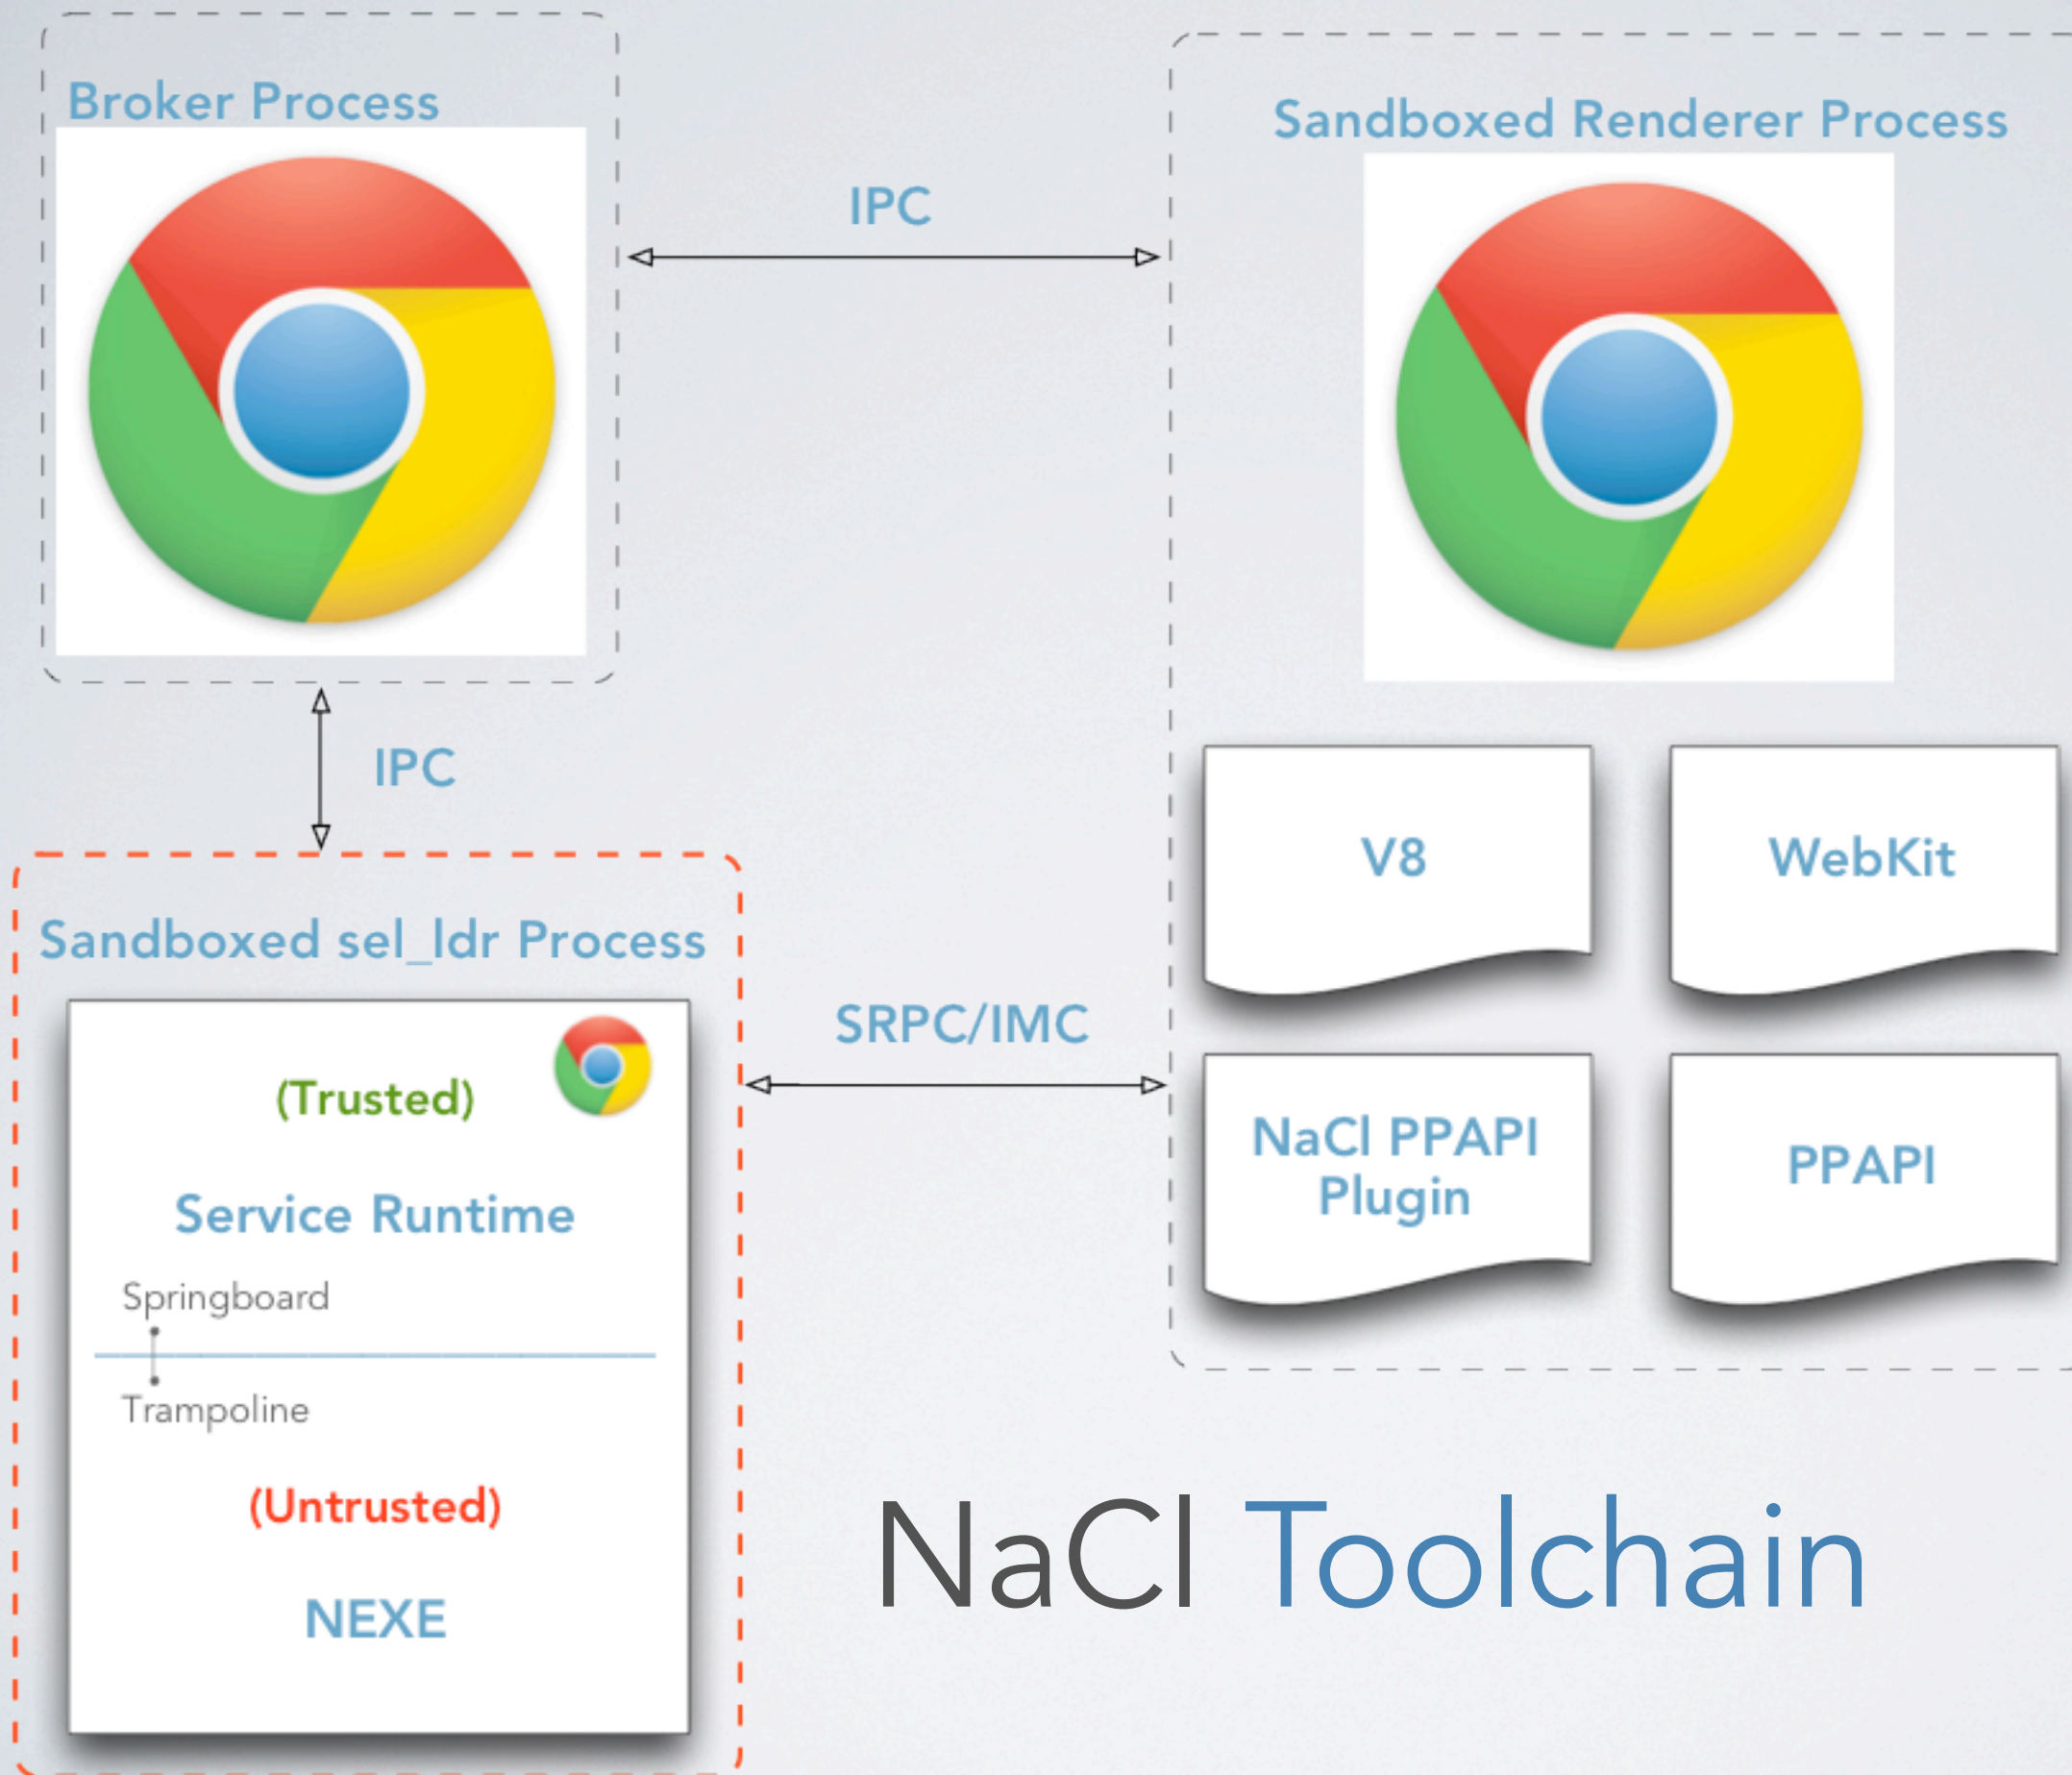
# Pepper Proxy

- Manages callbacks and related SRPC data structures
  - Callbacks exist to inform a remote NEXE when the renderer has completed an operation
  - Callbacks are bound to the SRPC channel
  - Callbacks can only be invoked on the main thread



# Pepper Proxy

- The pepper proxy is not a security boundary
- It performs little to no validation of untrusted data
- It exists only to proxy data between an untrusted NEXE and the PPAPI implementation



# NaCl Toolchain



# NaCl Toolchain

- Modified GCC toolchain ships with NaCl SDK
- Only the SDK compiler can be used to produce a NEXE
- SDK can produce code for:
  - 32-bit x86
  - 64-bit x86\_64



# NaCl Toolchain

- NEXE modules are compiled and linked as ELF
  - All of the typical structures are present
    - ELF Header, Program Headers, Dynamic Segment, Section Headers, Symbol Tables, Relocation Entries
- Standard tools can be used to disassemble a NEXE



# NaCl Toolchain

- NEXE instructions must be aligned to 32 byte boundary
  - This is required by the inner sandbox
- Blacklisted instructions are never emitted

01000ac0 <PpapiPluginStart>: ; 32 byte aligned

1000ac0:	53	push	%ebx
1000ac1:	83 ec 28	sub	\$0x28,%esp
1000ac4:	a1 00 00 02 11	mov	0x11020000,%eax
1000ac9:	8b 08	mov	(%eax),%ecx
1000acb:	85 c9	test	%ecx,%ecx
1000acd:	74 31	je	1000b00 <PpapiPluginStart+0x40>
1000acf:	eb 0f	jmp	1000ae0 <PpapiPluginStart+0x20>
1000ad1:	90	nop	
1000ad2:	90	nop	



# NaCl Toolchain

- No instructions may straddle the 32 byte boundary
- Branches are used to transfer control across boundaries
- No *ret* instructions, the stack is manually modified

01000ac0 <PpapiPluginStart>: ; 32 byte aligned

1000ac0:	53	push	%ebx
1000ac1:	83 ec 28	sub	\$0x28,%esp
1000ac4:	a1 00 00 02 11	mov	0x11020000,%eax
1000ac9:	8b 08	mov	(%eax),%ecx
1000acb:	85 c9	test	%ecx,%ecx
1000acd:	74 31	je	1000b00 <PpapiPluginStart+0x40>
1000acf:	eb 0f	jmp	1000ae0 <PpapiPluginStart+0x20>
1000ad1:	90	nop	
1000ad2:	90	nop	



# NaCl Toolchain

- Branch instructions are properly aligned to validated code
- call instructions are subject to a simple AND operation
- Alignment masking on the destination register ensures a 32 byte alignment which guarantees the destination has been run through the validator
- Prevents over written NEXE function pointers and modified registers from resulting in arbitrary code execution

```
0x100057b:    83 e0 e0    and    $0xfffffffffe0,%eax ; eax = 0x100057a
```

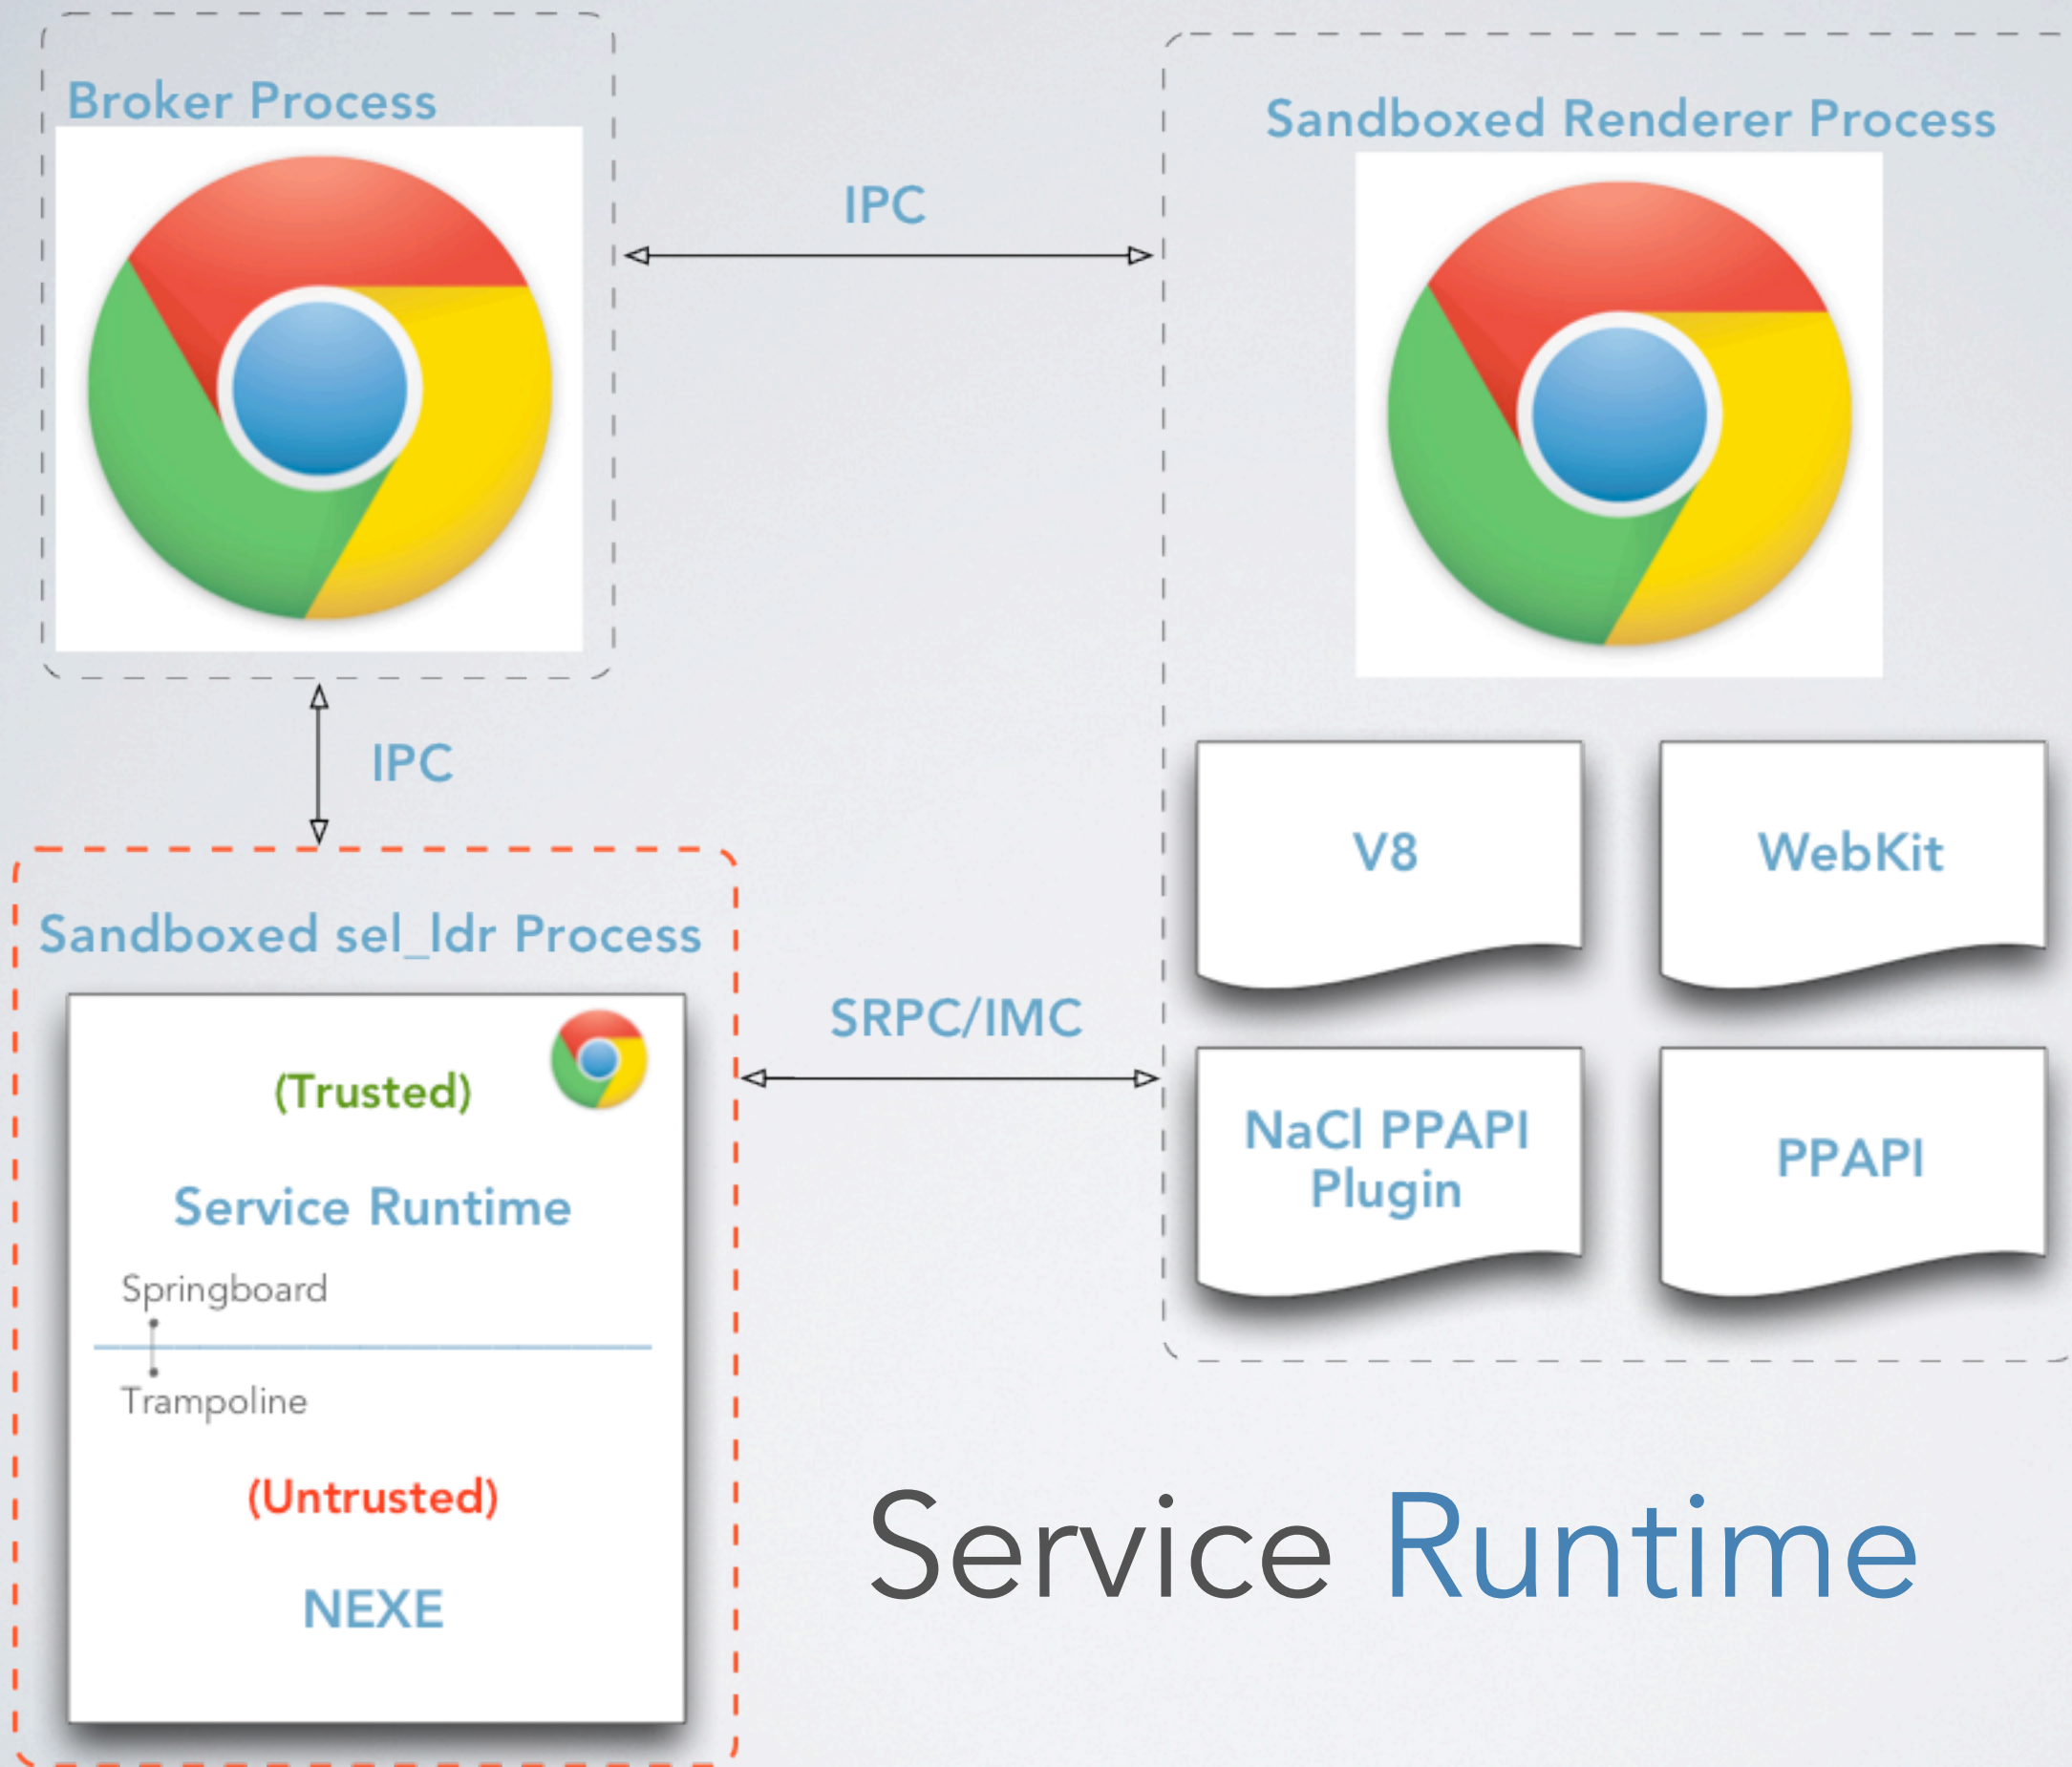
```
0x100057e:    ff d0      call   *%eax           ; eax = 0x1000560
```



# NaCl Toolchain

- NaCl toolchain
  - Ensures all ELF structures are properly formed
  - Only safe instructions are emitted
  - All branch instructions are preceded by alignment mask
- But an attacker can modify the NEXE binary...







# Service Runtime

- Stand alone process, launched from Chrome
- Runs in the same outer sandbox as Chrome renderer
- Shares its virtual address space with a NEXE module
  - Service Runtime memory is trusted
  - NEXE memory is untrusted
- Separation between the two is the 'Inner Sandbox'



# ELF Loader

- Multiple segments refer to other segments and so on
- Validates ELF header and Program headers
- ELF loader enforces:
  - Only one executable segment allowed
  - Segment load addresses and sizes



# Instruction Validator

- Disassembles all NEXE instructions
- Starts at trusted 32 byte aligned entry point
- Exits on any blacklisted instructions
  - Privileged instructions
  - Instructions that modify segment registers
  - ret
  - sysenter
  - prefix bytes



# Instruction Validator

- The validator only performs static analysis on the code
- Static analysis will not know register values at runtime
- As long as the branch target is mod 32 aligned it knows whatever the target is it has been validated
- Example:
  - 0x1000560 not known when the validator runs
    - Guaranteed to have been properly validated due to proper 32 byte alignment

```
0x100057b:    83 e0 e0    and    $0xffffffffe0,%eax ; eax = 0x100057a
```

```
0x100057e:    ff d0      call   *%eax           ; eax = 0x1000560
```



# Instruction Validator

- The NaCl SDK contains a standalone NEXE validator

```
$ ./ncval_x86_32 ../examples/hello_world_glibc/hw.nexe
segment[0] p_type 3 p_offset ee0 vaddr 11000ee0 paddr 11000ee0 align 1 filesz 19 memsz 19 flags 4
segment[1] p_type 1 p_offset 140 vaddr 1000140 paddr 1000140 align 65536 filesz da0 memsz da0 flags 5
parsing segment 1
```

**VALIDATOR: 10008e6: ret instruction (not allowed)**

**VALIDATOR: 10008e6: Illegal instruction**

```
segment[2] p_type 1 p_offset ee0 vaddr 11000ee0 paddr 11000ee0 align 65536 filesz 478 memsz 478 flags 4
segment[3] p_type 1 p_offset 1358 vaddr 11011358 paddr 11011358 align 65536 filesz 13c memsz 13c flags 6
segment[4] p_type 1 p_offset 10000 vaddr 11020000 paddr 11020000 align 65536 filesz 0 memsz 2c flags 6
segment[5] p_type 2 p_offset 136c vaddr 1101136c paddr 1101136c align 4 filesz d8 memsz d8 flags 6
segment[6] p_type 1685382481 p_offset 0 vaddr 0 paddr 0 align 4 filesz 0 memsz 0 flags 6
segment[7] p_type 7 p_offset 0 vaddr 0 paddr 0 align 4 filesz 0 memsz 0 flags 4
```

**\*\*\* ../examples/hello\_world\_glibc/hw.nexe IS UNSAFE \*\*\***

**Validated ../examples/hello\_world\_glibc/hw.nexe**

**\*\*\* ../examples/hello\_world\_glibc/hw.nexe IS UNSAFE \*\*\***



# Service Runtime

- Inner Sandbox
  - Enforced by loading the untrusted NEXE into the trusted service runtime address space
- Separating untrusted code and data
  - x86 memory segmentation model
  - x86\_64 mov/branch alignment, guard pages
  - ARM load/store/branch alignment, guard pages
  - Trusted TLS segments store data/registers between context switches



# Service Runtime

- x86 Segment Registers
  - Used to separate trusted from untrusted code/data
  - Modified when switching between trusted/untrusted
  - %cs code
  - %ds data
  - %gs thread local storage
  - %ss %es %fs all set to %ds
- Trusted instruction blocks are mapped at runtime to enable a context switch between trusted and untrusted
  - Springboard enable trusted to untrusted
  - Trampolines enable untrusted to trusted
  - Each contains privileged instructions that manipulate the segment registers



# Service Runtime

## Trusted Springboard

```
hlt  
mov 0x34(%ecx), %eax  
lss 0x1c(%ecx), %esp  
movw 0x28(%ecx), %ds  
jmp *%edx
```

## Untrusted NEXE Code

```
and 0xfffffffffe0, %ebx  
call *%ebx
```

The untrusted NEXE code is unable to transfer control to an arbitrary location in the trusted Springboard due to the alignment mask before the call. It may only execute the first instruction, which is a *hlt*



# Service Runtime

- No existing or new memory allocations may be marked as executable at runtime
- This guarantees only validated code pages have executable permissions



# NACL Syscalls

- A NEXE cannot make normal syscalls
  - sysenter / int 0x80 not allowed by the validator
- 30~ NACL\_SYSCALLS
- Basic operations such as open, close, read, write, ioctl, mmap, munmap, stat, \_exit and a few others
  - NACL specific IMC accept, connect, send, recv
- These are dispatched via Springboard/Trampoline code



# Service Runtime IMC

- Inter-Module Communication (IMC) is a protocol implemented by NaCl
- IMC uses basic socket types that are initialized by the `imc_makeboundsock` and `imc_socketpair` `NACL_SYSCALL`'s
- IMC is built on top of platform supplied UNIX sockets, named pipes and shared memory
- IMC is too low level to be used by application developers



# Service Runtime SRPC

- Simple Remote Procedure Call (SRPC)
- Is encapsulated by IMC
- SRPC allows for the encapsulation of serialized data between NEXE modules and the NaCl plugin
- SRPC endpoints are invoked via SRPC signatures

```
NaClSrpcInvokeBySignature(channel, "MyMethod:i:i", resource, bool)
```

- Although SRPC is at a higher level than IMC its not intended to be used by application developers directly



# Portable NaCl

- PNaCl (Pinnacle)
- Web site hosts LLVM IR produced by NaCl SDK toolchain
- PNaCl transforms LLVM IR to a NEXE for the users architecture using AOT compiler
- The inner sandbox and pepper proxy remain as-is
- 32-bit ARM Support
- Google has projected a 2012 release



# NaCl Attack Surface



**Aw, Snap!**

Something went wrong while displaying this webpage. To continue, reload or go to another page.

If you're seeing this frequently, try [these suggestions](#).



# NaCl Security

If a NaCl module can execute instructions that were not validated by the service runtime then the security provided by Native Client is broken



# NaCl Attack Surface

- Outer sandbox escape requires
  - Vulnerabilities in the broker process
  - No/Weakly sandboxed processes (Flash, GPU)
  - Kernel vulnerabilities the sandbox can't fully prevent
- Inner sandbox restrictions
  - Cannot reach the broker process directly
  - You cannot make syscalls or talk to the kernel directly
  - Instruction validation



# NaCl Attack Surface

- NaCl raises the bar for exploitation
- But you need trusted components to do useful things
  - These reside in the outer sandbox
- We find attack surface anywhere untrusted code can influence the execution of trusted code



# NaCl Attack Surface

- Vulnerable NEXE modules are not an issue
  - Validation of all direct/indirect execution branches
  - Why attack a NEXE module when you can load one directly instead?
- Malicious NEXE modules
  - Attempt to find and exploit vulnerabilities in various trusted NaCl components
  - **Find vulnerabilities in the Chrome renderer and go after them with an untrusted NEXE module**



# NaCl Attack Surface

- Service Runtime
  - Inner Sandbox
    - ELF loader
    - Instruction validator disassembler
    - NACL\_SYSCALL implementations



# NaCl Attack Surface

- NaCl PPAPI Plugin
  - IMC
  - SRPC
- JavaScript DOM interfaces
- JSON parser



# NaCl Attack Surface

- PPAPI - Pepper Proxy
  - Server interfaces
  - Client interfaces
- PPAPI
  - Interface implementations reached via the proxy
- GPU
  - Direct interfaces to the GPU
  - Pinky Pie exploited an integer overflow here at Pwnium



# First Inner Sandbox Breakout

- Call instruction memory dereference

```
andl $0xfffffffffe0, %edx  
call *(%edx)
```

- The validator and branch alignment ensure the value in the register is 32 byte aligned but not the value it references
- Results in execution of a non-validated instruction
- Discovered by Alex Radocea



# 2009 Security Contest

- Uncovered 20 new security vulnerabilities in NaCl
- Nothing that significantly broke the inner sandbox design
- 1st place - Mark Dowd, Ben Hawkes
  - 2nd place - Chris Rohlf, Eric Monti, Jason Carpenter
  - 3rd place - Gabriel Campana
  - 4th place - Daiki Fukumori
  - 5th place - Alex Radocea
- Some of these vulnerabilities are still relevant to NaCl



# 2009 Security Contest

- Two byte jmp prefixes
  - Unchecked prefix bytes on branch instruction
- EFLAGS direction flag modification
  - Untrusted NEXE code can change the direction flag
- Validated code unmapping
  - An untrusted NEXE can unmap valid code and replace it
- Uninitialized vtable
  - An uninitialized vtable can lead to code execution
- Double delete operator
  - NPAPI Object life cycle management confusion



# 2009 Security Contest

- The architecture has changed significantly since 2009
- NPAPI is no longer used but the same type of issues could be found in the PPAPI replacement
- Proved the inner sandbox was relatively strong
  - Provided a good look at the future of NaCl vulnerabilities
  - Trusted components that handle untrusted data are more likely to contain vulnerabilities than the inner sandbox



# Google NaCl Vulnerabilities

- Win64 inner sandbox escape via KiUserExceptionDispatcher
  - Exceptions transfer execution to this function which is not aware of the trusted vs untrusted stack
- bsf Instruction inner sandbox escape
  - Cannot properly validate alignment of conditional value
- Trampoline address space leak
  - Does not use PIC code, leaks a .text instruction address
- Address space leak via JavaScript error
  - A JavaScript error message contained a memory address



# Pepper Proxy Source Audit

- June 2011 Pepper Proxy Source Audit
  - Performed under contract to Google while at Matasano
- 3 week project
- 1 Person
- 10 vulnerabilities discovered
- Manual source audit of C/C++
  - Many thousands of lines
  - No scanners
  - Mostly grep and reading source



# Pepper Proxy Source Audit

- Are there vulnerabilities in the pepper proxy that would allow a malicious NEXE to escape the inner sandbox?
- Many interfaces accept and deserialize untrusted data
- Lots of opportunity for vulnerable code



# Pepper Proxy Source Audit

- PPB\_URLLoader\_Open CORS Header Injection

```
set_method = request_interface_->SetProperty(request_,  
    PP_URLREQUESTPROPERTY_METHOD,  
    Module::StrToVar("POST\x0d\x0ax-csrf-token:\x20test1234"));
```

- PPB\_Audio\_Create SRPC Channel Use After Free
  - The pepper proxy can be tricked into using a stale SRPC channel pointer when a callback is completed
  - Proof of concept:
    - Add `exit(-1)` after requesting PPAPI fill the audio buffer
    - Pepper proxy SRPC invalidates the channel but the callback structure is left with a stale pointer



# Pepper Proxy Source Audit

- PPB\_Graphics2D\_Create Shared Memory Overflow

- Check for overflow in (width \* height)

```
if (static_cast<int64>(width) * static_cast<int64>(height) >=
    std::numeric_limits<int32>::max())

    return false; // Prevent overflow of signed 32-bit ints.
```

- Perform multiplication after the check but \* 4

```
PepperPluginDelegateImpl::CreateImage2D(int width, int height) {
    uint32 buffer_size = width * height * 4;
}
```

- Create a new shared memory segment

```
TransportDIB::Create() {
    const int shmkey = shmget(IPC_PRIVATE, size, 0666);
}
```



# Pepper Proxy Source Audit

- Integer/Buffer Overflows

- PPB\_Graphics2D\_Create Shared Memory Integer Overflow
- PPB\_Context3DTrusted\_CreateTransferBuffer Shared Memory Integer Overflow
- PPB\_URLLoader\_ReadResponseBody Heap Overflow
- PPB\_FileIO\_Dev\_Read Heap Overflow
- PPB\_PDF\_SearchString Potential Heap Overflow
- MessageChannelEnumerate Heap Overflow

- Use After Free

- PPB\_Audio\_Create SRPC Channel Use After Free

- Information Leak

- PPB\_FileIO\_Write Out Of Bounds Read Information Leak

- Other

- PPB\_URLLoader\_Open CORS Request Allows For Header Injection
- PPB\_FileRef\_Create Potential Directory Traversal



# Pepper Proxy Source Audit

- Length calculations are difficult with serialized binary data
- Callbacks and channel related structures create complexities on top of application logic
- Confusion over whether PPAPI or the proxy should validate data can lead to vulnerabilities



# Chrome Shaker

- A pepper proxy fuzzer
- Developed under contract to Google January 2012
- Joint project with Matasano Security (Cody Brocious)
- Google deployed it in their fuzzing farm
- <https://code.google.com/p/chrome-shaker/>



# Chrome Shaker

- Simple NEXE template in C++
  - NEXE template code, random numbers, memory etc...
- Python tool that parses pepper proxy IDL files
  - Generates C++ into the basic template
  - Sets up each PPAPI interface in the 'C' style
  - Calls interfaces in random order with random arguments



# Chrome Shaker

- Fuzzing from a NEXE is not easy
- You can't log to disk
  - How do you write to disk from within two sandboxes?
    - Use the interface (FileIO) that you're currently fuzzing?
    - We had to resort to STDOUT
- You need a constant source of random data
  - We cheated and call into JavaScript for window.crypto
- Some code paths require calling interfaces in order
  - We have an API dependency file in yaml for this



# NaCl Exploitation

- Exploitation of a vulnerability in a trusted component allows arbitrary code execution in the outer sandbox
- Inner sandbox has been defeated
- Chrome renderer sandbox is still enforced
- This is not equivalent to a WebKit/V8 bug
  - The pepper proxy exposes interfaces at a lower level that JavaScript cannot reach



# The Conclusion

- NaCl is trying to solve a difficult problem
- NaCl is not ActiveX or NPAPI
- NaCl research into SFI and the inner sandbox will influence future sandbox designs with similar goals
- As the trusted code attack surface grows more implementation vulnerabilities will be found
  - The NaCl design helps to mitigate their impact
- Enable Chrome's Click-To-Play



# The End

## Questions?

leafsr.com  
Chris.Rohlf@gmail.com



# BlackHat Survey

Please fill out the BlackHat survey!