

Augmenting Binary Analysis with Python and Pin

November 14th, 2014

Who are we?

About Us

- Omar
 - Security Engineer at Etsy
 - Recent graduate of NYU
- Tyler
 - Studies at NYU

What is binary analysis?

What is binary analysis?

- *Binary*: A file containing all the resources and native code needed for a program to execute
- *Analysis*: To make sense of an application when the original intentions are not clear or known

Using a debugger (WinDbg, GDB, Immunity, etc)

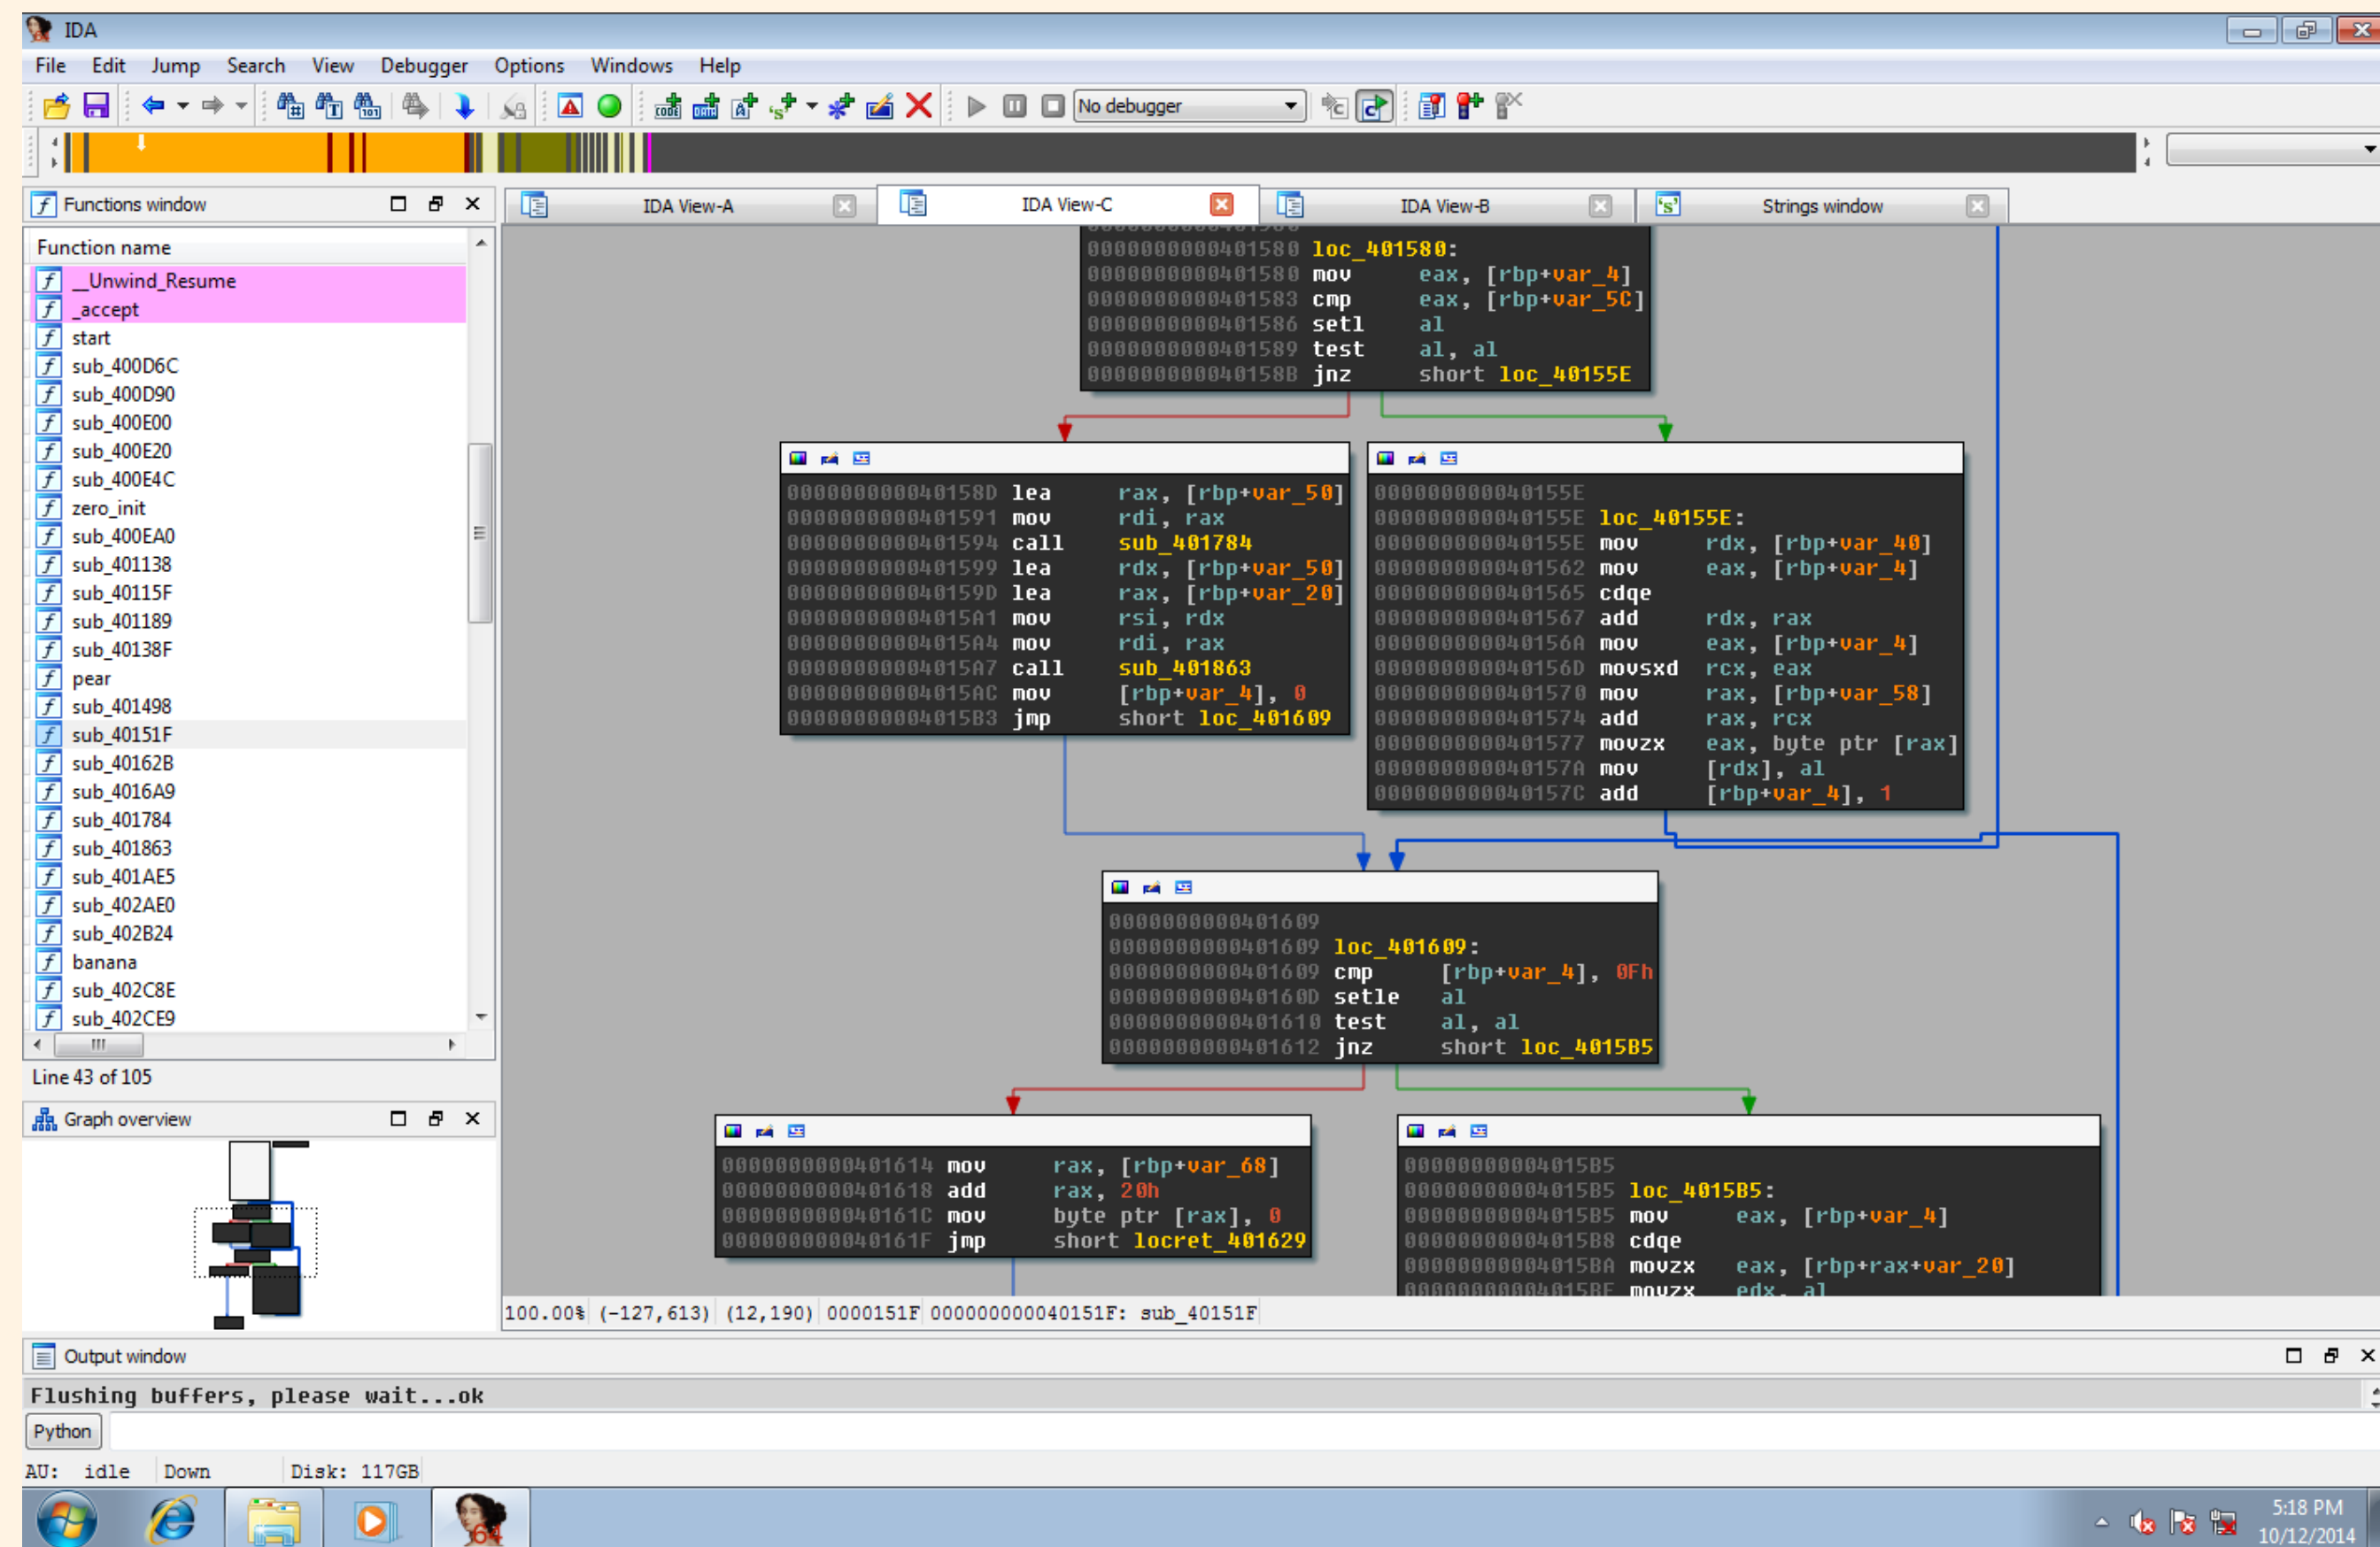
```
0x7ffff72771f3 <__GI___libc_malloc+35>:      mov     rbx,QWORD PTR fs:[rax]
(gdb) finish
Run till exit from #0  __GI___libc_malloc (bytes=140737488347216) at malloc.c:2876
0x000000000404128 in ?? ()
1: x/10i $pc
=> 0x404128:      mov     QWORD PTR [rbp-0x20],rax
0x40412c:      mov     eax,DWORD PTR [rip+0x20140e]      # 0x605540
0x404132:      mov     edx,DWORD PTR [rbp-0x14]
0x404135:      mov     edi,edx
0x404137:      sub     edi,eax
0x404139:      mov     eax,edi
0x40413b:      mov     edx,eax
0x40413d:      shr     edx,0x1f
0x404140:      add     eax,edx
0x404142:      sar     eax,1
Value returned is $3 = (void *) 0x6065f0
```

Simply observing the execution of a binary

```
$ ./bomb
Welcome to my fiendish little bomb. You have 6 phases with
which to blow yourself up. Have a nice day!
hello world

BOOM!!!
The bomb has blown up.
$ ./bomb
Welcome to my fiendish little bomb. You have 6 phases with
which to blow yourself up. Have a nice day!
Public speaking is very easy.
Phase 1 defused. How about the next one?
```

Reading disassembly output (IDA, objdump, etc)



Running /usr/bin/strings on a binary

```
So you think you can stop the bomb with ctrl-c, do you?
Well...
OK. :-)
Invalid phase%s
%d %d %d %d %d %d
Bad host (1).
Bad host (2).
Bad host (3).
Error: Premature EOF on stdin
GRADE_BOMB
Error: Input line too long
ERROR: dup(0) error
ERROR: close error
ERROR: tmpfile error
Subject: Bomb notification
nobody
defused
exploded
bomb-header:%s:%d:%s:%s:%d
bomb-string:%s:%d:%s:%d:%s
bomb
/usr/sbin/sendmail -bm
%s %s@%s
ERROR: notification error
ERROR: fclose(tmp) error
ERROR: dup(tmpstdin) error
ERROR: close(tmpstdin)
BOOM!!!
The bomb has blown up.
%d %s
austinpowers
Curses, you've found the secret phase!
But finding it and solving it are quite different...
Congratulations! You've defused the bomb!
generic
isrveawhobpnutfg
(
```

Static Analysis

- Reading disassembly output (IDA, objdump, etc)
- Running `/usr/bin/strings` on a binary
- Using a debugger (WinDbg, gdb, Immunity, etc)
- Simply observing the execution of a binary

Dynamic Analysis

- Reading disassembly output (IDA, objdump, etc)
- Running `/usr/bin/strings` on a binary
- Using a debugger (WinDbg, gdb, Immunity, etc)
- Simply observing the execution of a binary

Static vs Dynamic

- Speed
- Level of Understanding
- Code Coverage
 - Static can cover 100% of the code (good or bad?)
 - Dynamic can be accurate due to run time information

Introducing...

Dynamic Binary Instrumentation

Dynamic Binary Instrumentation

- A technique to modify the behavior of programs based on certain conditions during execution
- Generally, DBI involves redirecting execution to your code before or after the target executes
 - Sometimes done by modifying the code before starting the program
 - For example, an INT3 instruction on x86 used by debuggers, or less specifically, trampolines

Debugger Scripting

- GDB & LLDB
 - Scriptable using Python - Unix only (mostly)
- WinDBG
 - Scriptable using Python (somewhat) - Windows only
- VDB
 - Entirely Python API - Windows and Unix support
 - Lacks documentation, breaks often

Debugger Scripting

```
define structs
  set $target = $root
  set $limit = 0
  while $target
    printf "[0x%x] node.name=0x%x; node.value=0x%x; node.next=0x%x; node.prev=0x%x\n", $target, *($target), *($target+4), *($target+8), *($target+0xc)
    set $old_target = $target
    set $target = *($target+8)

    if $old_target == $target
      set $limit = $limit + 1
    end

    if $limit > 10
      printf "Infinite loop?\n"
      set $target = 0
    end
  end
end
end
```

DBI Frameworks

- Valgrind
 - GPL'd system for debugging and profiling Linux programs
 - Automatically detects many memory management and threading bugs
 - Works on x86/Linux, AMD64/Linux and PPC32/Linux
 - Focused on Safe and Reliable Code
 - Developer tool used for finding code errors

DBI Frameworks

- Valgrind
 - DEMO

DBI Frameworks

- Address Sanitizer
 - Fast memory error detector
 - The tool consists of a compiler instrumentation module (currently, an LLVM pass) and a run-time library which replaces the malloc function
 - Works on x86 Linux, and Mac, and ARM Android
 - Focused on bugs
 - Heap/Stack Buffer overflows and Use After Free

DBI Frameworks

- Address Sanitizer
 - DEMO????

DBI Frameworks

- DynamoRIO
 - Runtime code manipulation system that supports code transformations on any part of a program at runtime
 - Works on x86/AMD64 Linux Mac, and Windows
 - Transparent, and comprehensive manipulation of unmodified applications running on stock operating systems
 - Direct Competitor to Pin :-!

What *is* Pin?

- Pin allows user to insert arbitrary code into an executable right after it is loaded into memory
- Generates code from a “PinTool” used to “hook” instructions and calls
- Pin is the framework
- PinTools are the interface
 - The mechanism that decides where and what code is inserted
 - The code to execute at insertion points

Why Pin?

Intel's Pin

- Amazing documentation
- Same exact API works for Windows and Unix
- Extremely popular
- Nothing needs to be recompiled to be used with Pin

It's easy to get started

- Large repo of well commented sample tools come with Pin
- Documentation is generally easy to follow
- Installation is a piece of cake

It can be as granular as you need it to be

- Simple hook/callback system
 - instructions
 - basic blocks
 - function calls
 - and so on

Mostly personal preference, though

Why not Pin?

- The Pin API uses C++
 - Not a huge deal, but can be inconvenient during a time crunch (ctf)
 - Harder to prototype
- Slower than other DBI Frameworks
- Not as granular as other solutions
 - Harder to do more advanced binary analysis techniques such as taint tracing

Awesome but what can Pin do?

Popular Uses

- The Pin API has been used extensively in industry
- Most notably Microsoft Blue Hat (2012) Winner kBouncer (Vasilis Pappas)
 - Efficient and fully transparent ROP mitigation technique
 - Very similar to second place ROPGuard (Ivan Fratric)
 - Used in Microsofts EMET protection system
- IDA 6.4 and above includes a pin tool for tracing code in the debugger

Cool... WHERE ARE MY BUGS?!

- Pin can be used to find many different classes of bugs
- Most can be found by using the right kind of instrumentation
 - Format Strings
 - Analyze parameters passed to formatting functions
 - Buffer Overflows
 - Analyze memory read and write instructions
 - Misused Memory Allocation (Double Frees or UAF)
 - Analyze memory allocation functions (malloc/free) and memory writes

Misused Heap Allocations

- How to find these dynamically?
 - Keep track of all malloc calls and the addresses returned
 - Maintain state: Freed or In use and size
 - When a memory read or write happens, if the target is on the heap, verify that the memory is a valid place to be read from or written to

D-d-d-d-d-demo!

- Pin use after free demo

Pin

- Wow, Pin is really cool!
- But, wait! Pin is a mess!
 - Correction, C++ is a mess :P
- Lots of necessary boilerplate code
- Hard to prototype quickly
- Difficult to understand

C++

```
RTN mallocRtn = RTN_FindByName(img, MALLOC);
if (RTN_Valid(mallocRtn))
{
    RTN_Open(mallocRtn);

    // Instrument malloc() to print the input argument value and the return value.
    RTN_InsertCall(mallocRtn, IPOINT_BEFORE, (AFUNPTR) Arg1Before,
                  IARG_ADDRINT, MALLOC,
                  IARG_FUNCARG_ENTRYPOINT_VALUE, 0,
                  IARG_END);
    RTN_InsertCall(mallocRtn, IPOINT_AFTER, (AFUNPTR) MallocAfter,
                  IARG_FUNCRET_EXITPOINT_VALUE, IARG_END);

    RTN_Close(mallocRtn);
}
```

Python

```
rtn = pin.RTN_FindByName(img, "malloc")
if pin.RTN_Valid(rtn):
    pin.RTN_Open(rtn)
    pin.RTN_InsertCall(pin.IPOINT_BEFORE, "malloc", rtn, 1, malloc_before)
    pin.RTN_InsertCall(pin.IPOINT_AFTER, "malloc", rtn, 1, malloc_after)
    pin.RTN_Close(rtn)
```

C++ vs Python

- Python
 - Simpler
 - Cleaner
 - No need for recompilation every time
 - Extensive libraries and support

Python-Pin

- Essentially, a python interpreter embedded within a PinTool
 - “Virtual” pin module exposed to the python script
 - Enables access to most of Pin’s functionality from within python
 - Quick and easy to write PinTools
 - Negligible performance impact
 - Enables seamless integration with other Python modules
 - Z3py, PIL, SciPy, etc

Python-Pin Demo

- Use after free and double free
- Code coverage IDA basic block coloring
- Basic utility demos

The Future of Python-Pin

- Better memory management
- Implement more Pin functionality
- More testing for Windows support

Acknowledgements

Tyler Bohan
Kevin Chung
Dan Guido
Robert Meggs
Jonathan Salwan
Rich Smith
Paolo Soto
Alex Sotirov
Kai Zhong

Thanks for tuning in!

- Slides and pin tools will be posted to twitter, probably
 - @ancat/@1blankwall1