

# Transparent ROP Detection using CPU Performance Counters

他山之石，可以攻玉

Stones from other hills may serve to polish jade

**Xiaoning Li**

Intel Labs

**Michael Crouse**

Harvard University

THREADS Conference 2014

# Agenda

---

- Performance Monitor Architecture Overview
  - Past Research
  - Performance Events In Haswell/Silvermont
  - Mispredicted Branch Transfer
  - Stack Pivoting Detection Overview
  - Defense with CPU Performance Counters
  - APSA13-2 Case Study
  - Misc
-

---

# Performance Monitor Architecture Overview

---

# Performance Monitor Overview

---

- Introduced in the Pentium processor
    - A set of model-specific performance-monitoring counter MSRs
  - Enhanced to monitor a set of events in Intel P6 family of processors
  - Pentium 4 and Intel Xeon processors introduced a new performance monitoring mechanism and new set of performance events
  - Architectural/Non-Architectural performance events
-

# Architectural/Non-Architectural performance events

---

- The performance monitoring mechanisms and performance events defined for the Pentium, P6 family, Pentium 4, and Intel Xeon processors are not architectural
  - Intel Core Solo and Intel Core Duo processors support a set of architectural performance events and a set of non-architectural performance events
  - Processors based on Intel Core/Intel® Atom™ micro architecture support enhanced architectural performance events and non-architectural performance events
-

# Architectural/Non-Architectural performance events Availability

---

- Availability of architectural performance monitoring capabilities can be enumerated using the CPUID.0AH
- Non-architectural events for a given micro architecture can not be enumerated using CPUID

# CPUID Mechanism

---

- Number of performance monitoring counters available in a logical processor
  - Number of bits supported in each IA32\_PMCx
  - Number of architectural performance monitoring events supported in a logical processor
  - Software can use CPUID to discover architectural performance monitoring availability (CPUID.0AH).
-

# Architectural Performance Monitoring Leaf

---

<i>Architectural Performance Monitoring Leaf</i>		
0AH	EAX	Bits 07 - 00: Version ID of architectural performance monitoring Bits 15- 08: Number of general-purpose performance monitoring counter per logical processor Bits 23 - 16: Bit width of general-purpose, performance monitoring counter Bits 31 - 24: Length of EBX bit vector to enumerate architectural performance monitoring events
	EBX	Bit 00: Core cycle event not available if 1 Bit 01: Instruction retired event not available if 1 Bit 02: Reference cycles event not available if 1 Bit 03: Last-level cache reference event not available if 1 Bit 04: Last-level cache misses event not available if 1 Bit 05: Branch instruction retired event not available if 1 Bit 06: Branch mispredict retired event not available if 1 Bits 31- 07: Reserved = 0
	ECX	Reserved = 0
	EDX	Bits 04 - 00: Number of fixed-function performance counters (if Version ID > 1) Bits 12- 05: Bit width of fixed-function performance counters (if Version ID > 1) Reserved = 0



# Architectural Performance Monitoring Versions

---

CPU Architecture	Version
Intel Core Solo	1
Intel Core Duo	1
Core 2 Duo processor T7700	2
Core	2
Atom	2,3
Core i7	2,3

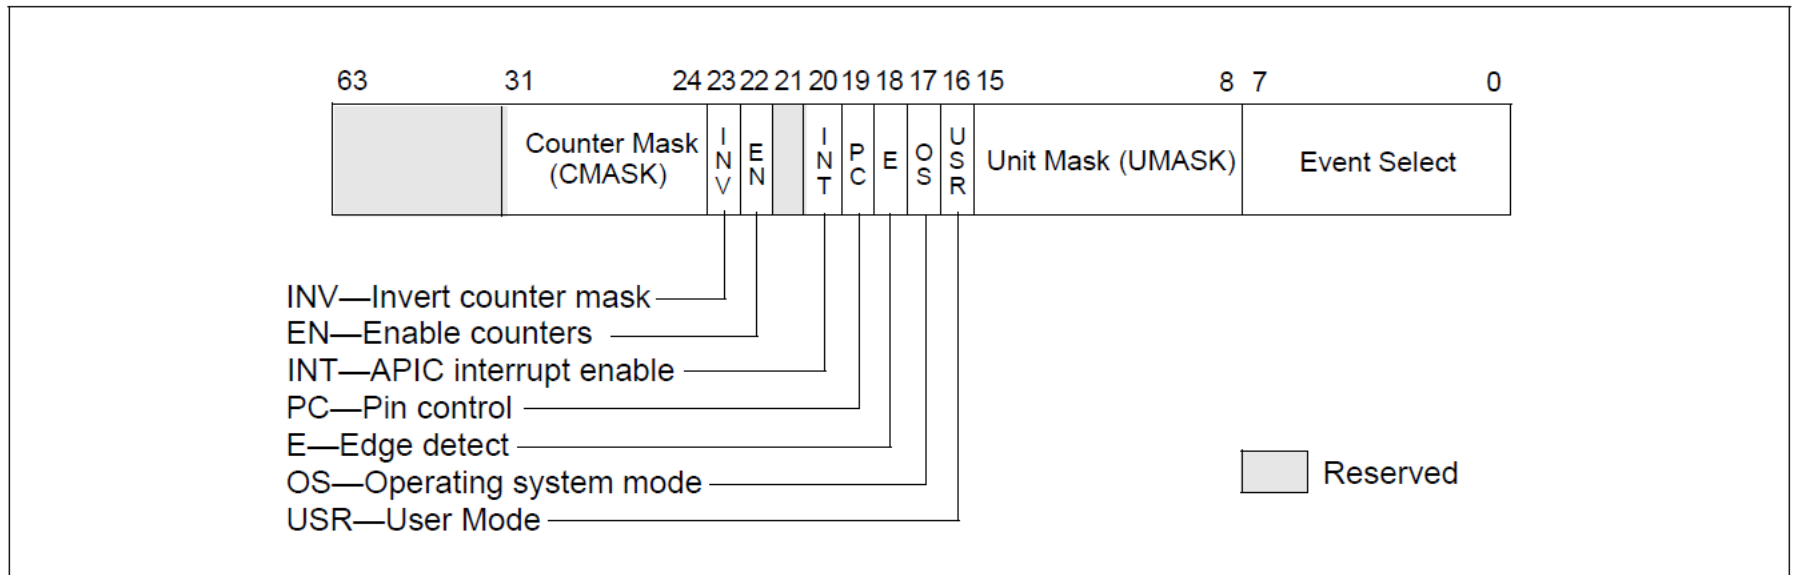
---

# Architectural Performance Monitoring Version 1

---

- Configure an architectural performance monitoring event with programming performance event select registers.
  - Performance event select MSRs (IA32\_PERFECTSELx MSRs).
  - Performance monitoring counter (IA32\_PMCx MSR)
  - Performance monitoring counters are paired with performance monitoring select registers.
-

# Layout of IA32\_PERFEVTSELx MSRs



# IA32\_PERFVTSELx MSRs

---

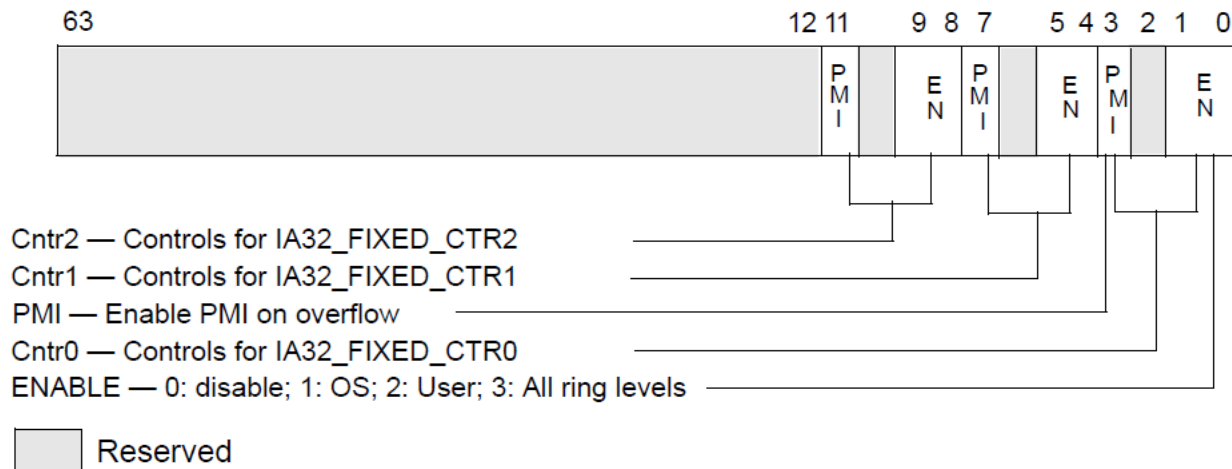
- Event select field (bits 0 through 7)
  - Unit mask (UMASK) field (bits 8 through 15)
  - USR (user mode) flag (bit 16)
  - OS (operating system mode) flag (bit 17)
  - E (edge detect) flag (bit 18)
  - PC (pin control) flag (bit 19)
  - INT (APIC interrupt enable) flag (bit 20)
  - EN (Enable Counters) Flag (bit 22)
  - INV (invert) flag (bit 23)
  - Counter mask (CMASK) field (bits 24 through 31)
-

# Architectural Performance Monitoring Version 2

---

- Bits 0 through 4 of CPUID.0AH.EDX indicates the number of fixed-function performance counters available per core
- Bits 5 through 12 of CPUID.0AH.EDX indicates the bit-width of fixed-function performance counters. Bits beyond the width of the fixed-function counter are reserved and must be written as zeros.


# IA32\_FIXED\_CTR\_CTRL MSR




Event Name	Fixed-Function PMC	PMC Address
INST_RETIRED.ANY	MSR_PERF_FIXED_CTR0/IA32_FIXED_CTR0	309H
CPU_CLK_UNHALTED.CORE	MSR_PERF_FIXED_CTR1//IA32_FIXED_CTR1	30AH
CPU_CLK_UNHALTED.REF	MSR_PERF_FIXED_CTR2//IA32_FIXED_CTR2	30BH

# IA32\_PERF\_GLOBAL\_CTRL MSR

---




# IA32\_PERF\_GLOBAL\_STATUS MSR





# IA32\_PERF\_GLOBAL\_OVF\_CTRL MSR

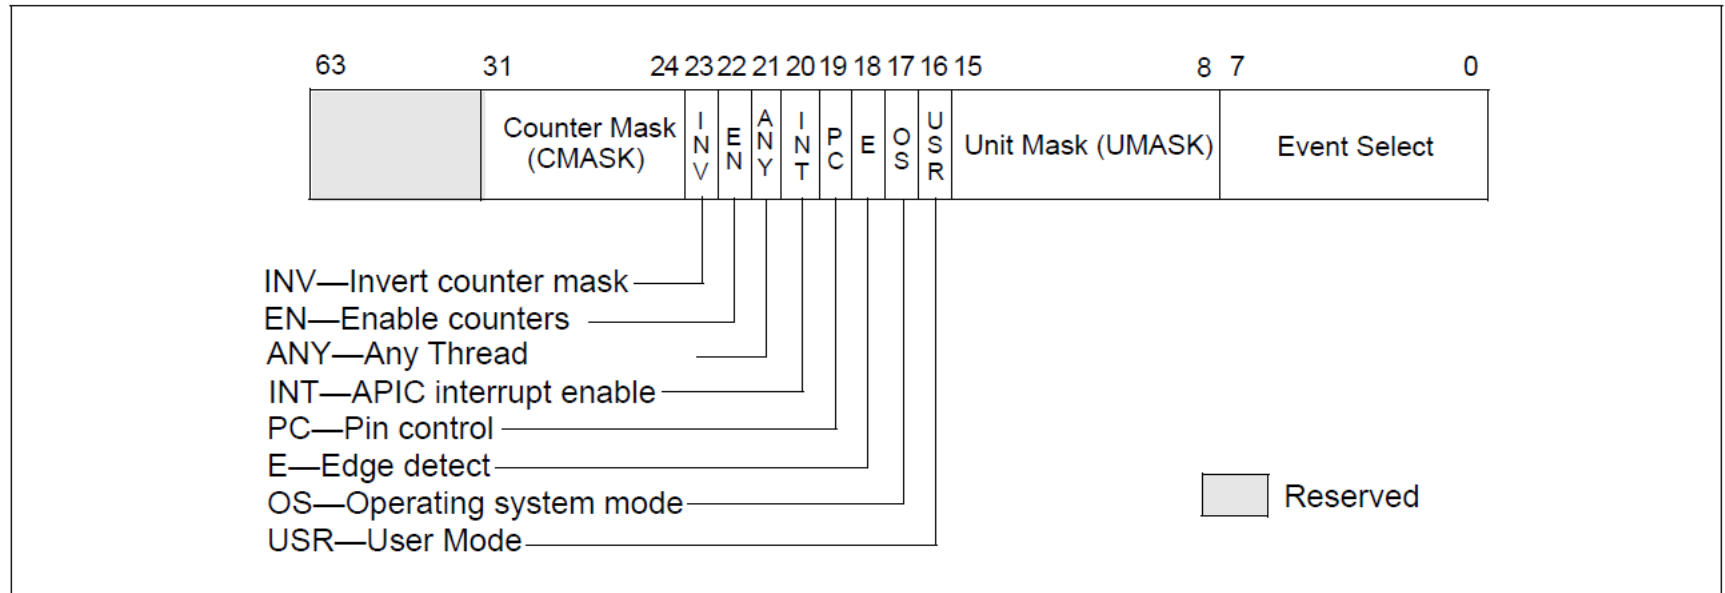


# Architectural Performance Monitoring Version 3

---

- The number of general-purpose performance counters (IA32\_PMCx) is reported in CPUID.0AH:EAX[15:8], the bit width of general-purpose performance counters is reported in CPUID.0AH:EAX[23:16]
- The bit vector representing the set of architectural performance monitoring events supported

# MSRs Supporting Architectural Performance Monitoring Version 3

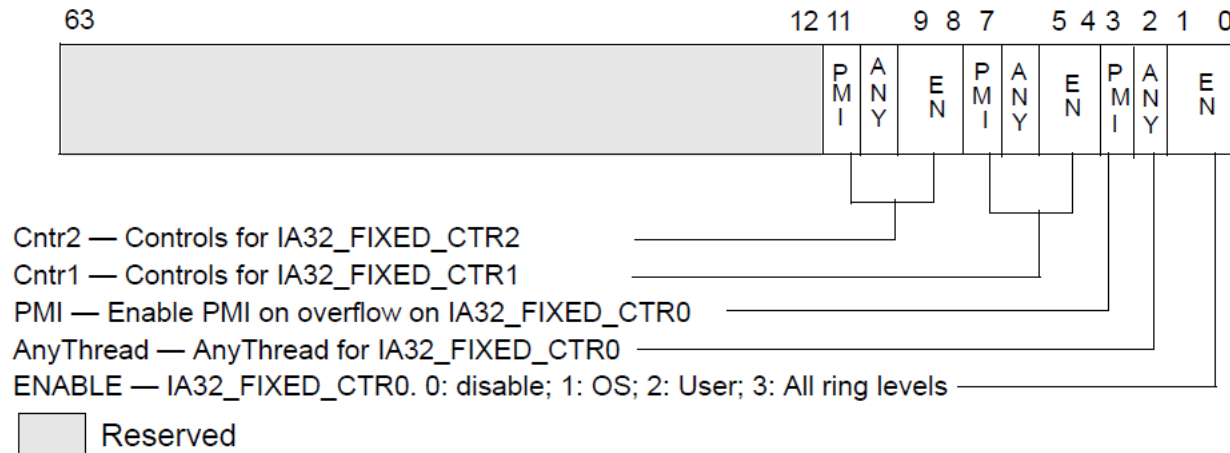


# IA32\_PERFEVTSELx MSRs Supporting Architectural Performance Monitoring Version 3

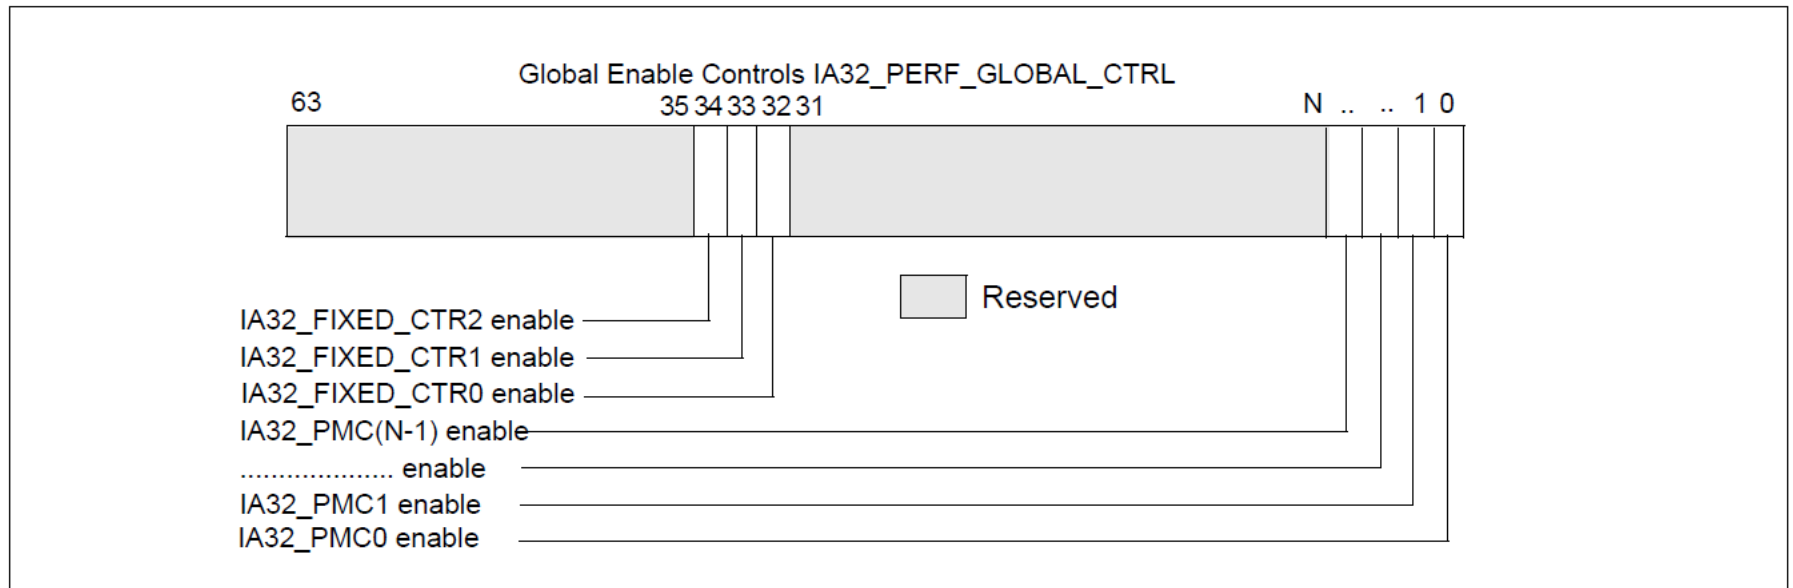
---

- Bit 21 (AnyThread) of IA32\_PERFEVTSELx is supported in architectural performance monitoring version 3.
    - When set to 1, it enables counting the associated event conditions (including matching the thread's CPL with the OS/USR setting of IA32\_PERFEVTSELx) occurring across all logical processors sharing a processor core.
    - When bit 21 is 0, the counter only increments the associated event conditions (including matching the thread's CPL with the OS/USR setting of IA32\_PERFEVTSELx) occurring in the logical processor which programmed the IA32\_PERFEVTSELx MSR.
-

# IA32\_FIXED\_CTR\_CTRL MSR Supporting Architectural Performance Monitoring Version 3




# IA32\_PERF\_GLOBAL\_CTRL MSR Supporting Architectural Performance Monitoring Version 3



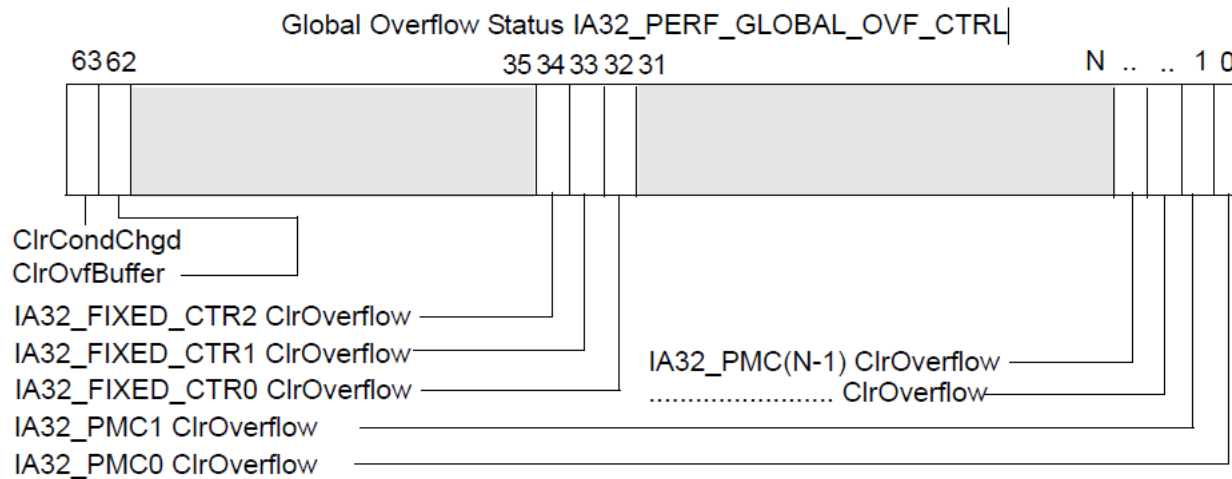
# IA32\_PERF\_GLOBAL\_STATUS MSR Supporting Architectural Performance Monitoring Version 3

---



# IA32\_PERF\_GLOBAL\_OVF\_CTRL MSR Supporting Architectural Performance Monitoring Version 3

---





# Pre-defined Architectural Performance Events

---

Bit Position CPUID.AH.EBX	Event Name	UMask	Event Select
0	UnHalted Core Cycles	00H	3CH
1	Instruction Retired	00H	C0H
2	UnHalted Reference Cycles	01H	3CH
3	LLC Reference	4FH	2EH
4	LLC Misses	41H	2EH
5	Branch Instruction Retired	00H	C4H
6	Branch Misses Retired	00H	C5H

---

# Branch Instructions Retired Events

---

- Branch Instructions Retired — Event select C4H, Umask 00H This event counts branch instructions at retirement. It counts the retirement of the last micro-op of a branch instruction.
- All Branch Mispredict Retired — Event select C5H, Umask 00H

---

# Past Research

---

## Security Breaches as PMU Deviation: Detecting and Identifying Security Attacks Using Performance Counters in **APSys'11, 2011**

---

- On Linux 2.6.34
- Using Machine Learning with performance counters
- Using PEBS/BTS for branch record
- Mentioned Branch Miss Predict Event

Source: [http://ipads.se.sjtu.edu.cn/\\_media/publications:perf-apsys.pdf](http://ipads.se.sjtu.edu.cn/_media/publications:perf-apsys.pdf)

Attack Type	Description	PMU Events
Code-injection	Inject code and take control transfer to injected code	Branch Tracing Event(BTS)
		Branch Miss Predict Event
		Instruction TLB Misses Event
Return-to-libc	Use library calls instead of inject code (e.g., invoke “execve” with “bin/bash”)	Branch Tracing Event(BTS)
Return-oriented programming	Use instructions before “ret” in existing library and binary code to form shellcode	Branch Tracing Event(BTS)

**Table 1: Deviation in performance characteristics of common attacks.**

---

## CFIMon: Detecting violation of control flow integrity using performance counters in **Dependable Systems and Networks (DSN), 2012**

---

- On Linux
- CFI defense with PMU events
- Using PEBS/BTS for branch record
- Target all branch
- Using call\_set/Ret\_set policy

Source: [http://ipads.se.sjtu.edu.cn/\\_media/publications:cfimon.pdf](http://ipads.se.sjtu.edu.cn/_media/publications:cfimon.pdf)

Branch Type	Branch Example	Target Instruction	Target Set	In Binary	Run-time
Direct call	callq 34df0 <abort>	1: taken	/	16.8%	14.5%
Direct jump	jnz c2ef0 <__write>	1 or 2: taken or fallthrough	/	74.3%	0.8%
Return	retq	Limited: insn. next to a call	<i>ret_set</i>	6.3%	16.3%
Indirect call	callq *%rax	Limited: 1st insn. of a function	<i>call_set</i>	2.1%	0.2%
Indirect jump	jmpq *%rdx	Unlimited: potentially any insn.	<i>train_set</i>	0.5%	68.3%

TABLE 1. Branch Classification. The distribution is from Apache and libraries it uses.

# Taming the ROPE on Sandy Bridge in SYSCAN 2013

---

- On Linux to detect Ring 0 ROP on SandyBridge
- Use BR\_MISP\_EXEC to catch RET misprediction event to detect ROP
- Using preceding call policy

Source: Taming the ROPE on Sandy Bridge in SYSCAN 2013

- 0x89 – BR\_MISP\_EXEC.\*: mispredicted executed branches
  - 0x800 – .RETURN\_NEAR: normal, near ret
  - 0x8000 – .TAKEN: unconditional branch
-

# HDROP: Detecting ROP Attacks Using Performance Monitoring Counters in **ISPEC'14**

---

- On Linux
- Use BR\_RET\_MISSP\_EXEC Event , which is replaced by 0x89
- Expect BR\_RET\_MISSP\_RETIRED, but not existed

---

# Performance Events in Haswell/Silvermont

---



# BR\_INST\_EXEC Event and UMask

---

88H	01H	BR_INST_EXEC.COND	Qualify conditional near branch instructions executed, but not necessarily retired.	Must combine with umask 40H, 80H
88H	02H	BR_INST_EXEC.DIRECT_JMP	Qualify all unconditional near branch instructions excluding calls and indirect branches.	Must combine with umask 80H
88H	04H	BR_INST_EXEC.INDIRECT_JMP_NON_CALL_RET	Qualify executed indirect near branch instructions that are not calls nor returns.	Must combine with umask 80H
88H	08H	BR_INST_EXEC.RETURN_NEAR	Qualify indirect near branches that have a return mnemonic.	Must combine with umask 80H
88H	10H	BR_INST_EXEC.DIRECT_NEAR_CALL	Qualify unconditional near call branch instructions, excluding non call branch, executed.	Must combine with umask 80H
88H	20H	BR_INST_EXEC.INDIRECT_NEAR_CALL	Qualify indirect near calls, including both register and memory indirect, executed.	Must combine with umask 80H
88H	40H	BR_INST_EXEC.NONTAKEN	Qualify non-taken near branches executed.	Applicable to umask 01H only
88H	80H	BR_INST_EXEC.TAKEN	Qualify taken near branches executed. Must combine with 01H, 02H, 04H, 08H, 10H, 20H.	
88H	FFH	BR_INST_EXEC.ALL_BRANCHES	Counts all near executed branches (not necessarily retired).	

---

# BR\_MISP\_EXEC Event and UMask

---

89H	01H	BR_MISP_EXEC.COND	Qualify conditional near branch instructions mispredicted.	Must combine with umask 40H, 80H
89H	04H	BR_MISP_EXEC.INDIRECT_JMP_NON_CALL_RET	Qualify mispredicted indirect near branch instructions that are not calls nor returns.	Must combine with umask 80H
89H	08H	BR_MISP_EXEC.RETURN_NEAR	Qualify mispredicted indirect near branches that have a return mnemonic.	Must combine with umask 80H
89H	10H	BR_MISP_EXEC.DIRECT_NEAR_CALL	Qualify mispredicted unconditional near call branch instructions, excluding non call branch, executed.	Must combine with umask 80H
89H	20H	BR_MISP_EXEC.INDIRECT_NEAR_CALL	Qualify mispredicted indirect near calls, including both register and memory indirect, executed.	Must combine with umask 80H
89H	40H	BR_MISP_EXEC.NONTAKEN	Qualify mispredicted non-taken near branches executed.	Applicable to umask 01H only
89H	80H	BR_MISP_EXEC.TAKEN	Qualify mispredicted taken near branches executed. Must combine with 01H,02H, 04H, 08H, 10H, 20H.	
89H	FFH	BR_MISP_EXEC.ALL_BRANCHES	Counts all near executed branches (not necessarily retired).	

# BR\_INST\_RETIRED Event and UMask

---

C4H	00H	BR_INST_RETIRED.ALL_BRANC HES	Branch instructions at retirement.	See <b>Table 19-1</b>
C4H	01H	BR_INST_RETIRED.CONDITIONA L	Counts the number of conditional branch instructions retired.	Supports PEBS
C4H	02H	BR_INST_RETIRED.NEAR_CALL	Direct and indirect near call instructions retired.	Supports PEBS
C4H	04H	BR_INST_RETIRED.ALL_BRANC HES	Counts the number of branch instructions retired.	Supports PEBS
C4H	08H	BR_INST_RETIRED.NEAR_RETU RN	Counts the number of near return instructions retired.	Supports PEBS
C4H	10H	BR_INST_RETIRED.NOT_TAKEN	Counts the number of not taken branch instructions retired.	
C4H	20H	BR_INST_RETIRED.NEAR_TAKE N	Number of near taken branches retired.	Supports PEBS
C4H	40H	BR_INST_RETIRED.FAR_BRANC H	Number of far branches retired.	

---

# BR\_MISP\_RETIREDD Event and UMask

---

C5H	00H	BR_MISP_RETIREDD.ALL_BRANC HES	Mispredicted branch instructions at retirement	See Table 19-1
C5H	01H	BR_MISP_RETIREDD.CONDITIONA L	Mispredicted conditional branch instructions retired.	Supports PEBS
C5H	04H	BR_MISP_RETIREDD.ALL_BRANC HES	Mispredicted macro branch instructions retired.	Supports PEBS
C5H	20H	BR_MISP_RETIREDD.NEAR_TAKE N	Number of near branch instructions retired that were taken but mispredicted.	

# BR\_MISP\_RETIRE In Silvermont

C5H	00H	BR_MISP_RETIRE.ALL_BRANCHES	Retired mispredicted branch instructions	This event counts the number of mispredicted branch instructions retired.
C5H	7EH	BR_MISP_RETIRE.JCC	Retired mispredicted conditional jumps	This event counts the number of mispredicted branch instructions retired that were conditional jumps.
C5H	BFH	BR_MISP_RETIRE.FAR	Retired mispredicted far branch instructions	This event counts the number of mispredicted far branch instructions retired.
C5H	EBH	BR_MISP_RETIRE.NON_RETURN_IND	Retired mispredicted instructions of near indirect jmp or call	This event counts the number of mispredicted branch instructions retired that were near indirect call or near indirect jmp.
C5H	F7H	BR_MISP_RETIRE.RETURN	Retired mispredicted near return instructions	This event counts the number of mispredicted near RET branch instructions retired
C5H	F9H	BR_MISP_RETIRE.CALL	Retired mispredicted near call instructions	This event counts the number of mispredicted near CALL branch instructions retired
C5H	FBH	BR_MISP_RETIRE.IND_CALL	Retired mispredicted near indirect call instructions	This event counts the number of mispredicted near indirect CALL branch instructions retired
C5H	FDH	BR_MISP_RETIRE.REL_CALL	Retired mispredicted near relative call instructions	This event counts the number of mispredicted near relative CALL branch instructions retired
C5H	FEH	BR_MISP_RETIRE.TAKEN_JCC	Retired mispredicted conditional jumps that were predicted taken	This event counts the number of mispredicted branch instructions retired that were conditional jumps and predicted taken.

---

# Mispredicted Branch Transfer and Exploits

---

---

Most non-logic vulnerabilities/exploits will cause  
**unexpected code flows!**

---

---

**Branch Misprediction Unit** is designed very well

---



---

**Branch Misprediction Events** based approach  
will catch enough events caused by exploits with  
reasonable performance impact

---

# Exploit Code And Misprediction

---

First unintended/intended exploit code triggered by **indirect CALL**



Indirect CALL misprediction

First unintended/intended exploit code triggered by **indirect JMP**



Indirect JMP misprediction

First unintended/intended exploit code triggered by **RET**



RET misprediction

Following unintended/intended ROP gadgets



Indirect CALL/JMP, RET misprediction

---

# Stack Pivoting Detection Overview

---

# Stack Pivoting Detection Solution

---

- Stack Pivoting needs to point stack pointer to customized data buffer, usually it's from heap
- Current detection solution
  - Critical APIs check

# Stack Pivoting Detection

---

- Validate stack limitation
- From FS:18h
- Use `_NT_TIB` or `_TEB` to get stack limitation

```
0:000> dt _NT_TIB
ntdll!_NT_TIB
+0x000 ExceptionList      : Ptr32 _EXCEPTION_REGISTRATION_RECORD
+0x004 StackBase          : Ptr32 Void
+0x008 StackLimit         : Ptr32 Void
+0x00c SubSystemTib       : Ptr32 Void
+0x010 FiberData          : Ptr32 Void
+0x010 Version            : Uint4B
+0x014 ArbitraryUserPointer : Ptr32 Void
+0x018 Self               : Ptr32 _NT_TIB
```

# Hooked API Examples IN EMET

---

## ➤ kernel32![API]Stub

### ➤ Hooked APIs

- kernel32.MapViewOfFileEx
  - kernel32.MapViewOfFile
  - kernel32.CreateFileMappingW
  - kernel32.CreateFileMappingA
  - kernel32.CreateFileW
  - kernel32.CreateFileA
  - kernel32.WinExec
  - kernel32.WriteProcessMemory
  - kernel32.CreateRemoteThread
  - kernel32.CreateProcessInternalW
  - kernel32.CreateProcessInternalA
  - kernel32.CreateProcessW
  - kernel32.CreateProcessA
  - kernel32.HeapCreate
  - kernel32.VirtualAllocEx
  - kernel32.VirtualAlloc
  - kernel32.LoadLibraryExW
  - kernel32.LoadLibraryExA
  - kernel32.LoadLibraryW
  - kernel32.LoadLibraryA
  - kernel32.VirtualProtectEx
  - kernel32.VirtualProtect
-

# Hooks Code Example

---

```
kernel32!VirtualProtectStub:  
76692bcd e93ed595f9 jmp 6fff0110  
76692bd2 5d pop ebp  
76692bd3 e900f5fbff jmp kernel32!VirtualProtect (766520d8)  
.....
```


```
no prior disassembly possible  
6fff0110 68516c38d4 push 0D4386C51h  
6fff0115 60 pushad  
6fff0116 9c pushfd  
6fff0117 54 push esp  
6fff0118 e893d399fd call emet+0x4d4b0 (6d98d4b0)  
6fff011d 9d popfd  
6fff011e 61 popad  
6fff011f 83c404 add esp,4  
6fff0122 8bff mov edi,edi  
6fff0124 55 push ebp  
6fff0125 8bec mov ebp,esp  
6fff0127 e9a62a6a06 jmp kernel32!VirtualProtectStub+0x5 (76692bd2)  
6fff0128 -- jmp ?
```

---

# StackPivot in EMET

---

```
GetStack_Limits(&stacklimit, &stackbase);  
if ( caller_esp < stacklimit || caller_esp > stackbase )  
{
```



```
}
```

---



# Improved Stack Pivoting Detection

---

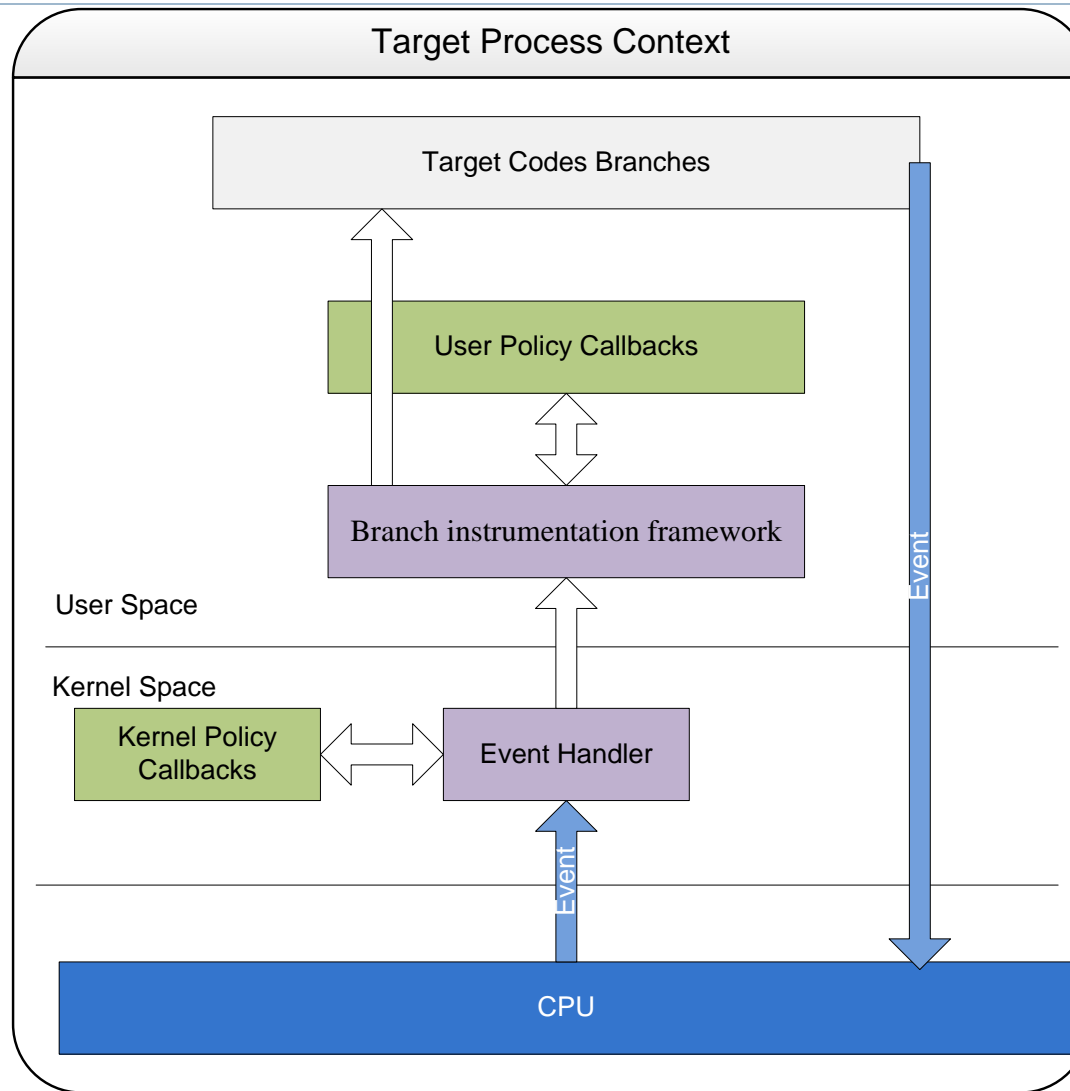
- Problems in API based approach
    - Hook hopping could bypass the check
    - Valid stack pointer before API calling will bypass the check
  - Improvement
    - Check on more APIs
    - Check on instruction
    - Check on branch instruction
    - Check on CALL/JMP branch instruction
    - Check on indirect CALL/JMP branch instruction
    - Check on mispredicted Indirect CALL/JMP branch instruction
-

---

# Defense with CPU Performance Counters

---

# Defense with CPU Performance Counters



# Defense on Windows 7 32bits

---

- MSR Programming
  - Counter configuration
  - Register Event Handler
  - Interrupt Context on Stack
  - Interrupt Event Handler
-

# Write MSR

---

## ➤ Compiler Intrinsics

`void __writemsr( unsigned long Register, unsigned __int64 Value );`

## ➤ Inline ASM

```
mov ecx, MSRID
mov edx, HIGH32b
mov eax, LOW32b
;wrmsr
__emit 00Fh
__emit 030h
```

### WRMSR—Write to Model Specific Register

Opcode	Instruction	Op/ En	64-Bit Mode	Compat/ Leg Mode	Description
0F 30	WRMSR	NP	Valid	Valid	Write the value in EDX:EAX to MSR specified by ECX.

# MSR Programming

---

- Use PMC0 to trigger PMI
  1. IA32\_DEBUGCTL
  2. LBR\_SELECT
  3. MSR\_PERF\_GLOBAL\_CTRL
  4. PMC0
  5. IA32\_PERF\_EVTSEL0
  6. MSR\_PERF\_GLOBAL\_CTRL

# Counter configuration

---

- Trigger Interrupt for every event
  - Set PMC0 = 0xffffffffffffff
- Interrupt will happen for every event cross whole system

# Register Event Handler

---

- Interrupt Vector - 0xFE
  - Set up event handler
    - IDT Hook
      - Replace vector 0xFE handler
    - API Hook
      - Hal!HalpPerfInterrupt
    - Callback Hook

```
mov     ecx, ebp
mov     eax, ds:_HalpPerfInterruptHandler
or      eax, eax
jz      short loc_8002B354
call    eax
```
  - Register via API
-



# Interrupt Context on Stack

---

- IDT Hook needs to maintain interrupt context as usual
- Callback Hook needs to address undocumented interrupt context from stack

```
typedef struct {  
    UINT32 stacktop;  
    UINT32 intno;  
    UINT32 esp;  
    UINT32 eip;  
    UINT32 Data[24];  
    UINT32 int_eip;  
    UINT32 int_cs;  
    UINT32 int_eflag;  
    UINT32 int_esp;  
    UINT32 int_ss;  
} IA32_PERFINTERRUPT_PARAMETER;
```

---

# Interrupt Event Handler

---

- Check IA32\_PERF\_GLOBAL\_STATUS and clear MSR\_PERF\_GLOBAL\_OVF\_CTL if multiple PMCs are used
  - Check CS in interrupt frame to filter out all ring 0 events
  - Check current CR3 is targeted process
  - Carefully deal with pagable memory to get stack range, code@From, code@To
    - APC, IRQL changes
  - Compare stack in interrupt frame with TIB stack range
  - Get last branch transfer record from LBR TOS
  - Clear IA32\_PERF\_GLOBAL\_OVF\_CTRL if need
-

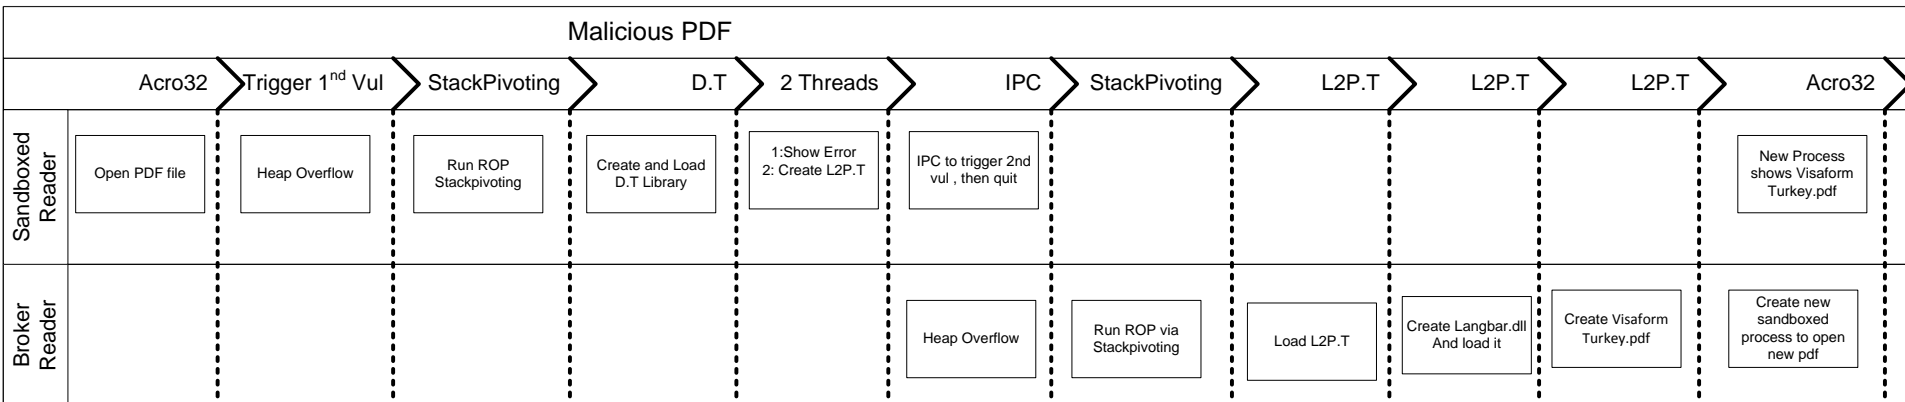
---

# APSA13-2 Case Study

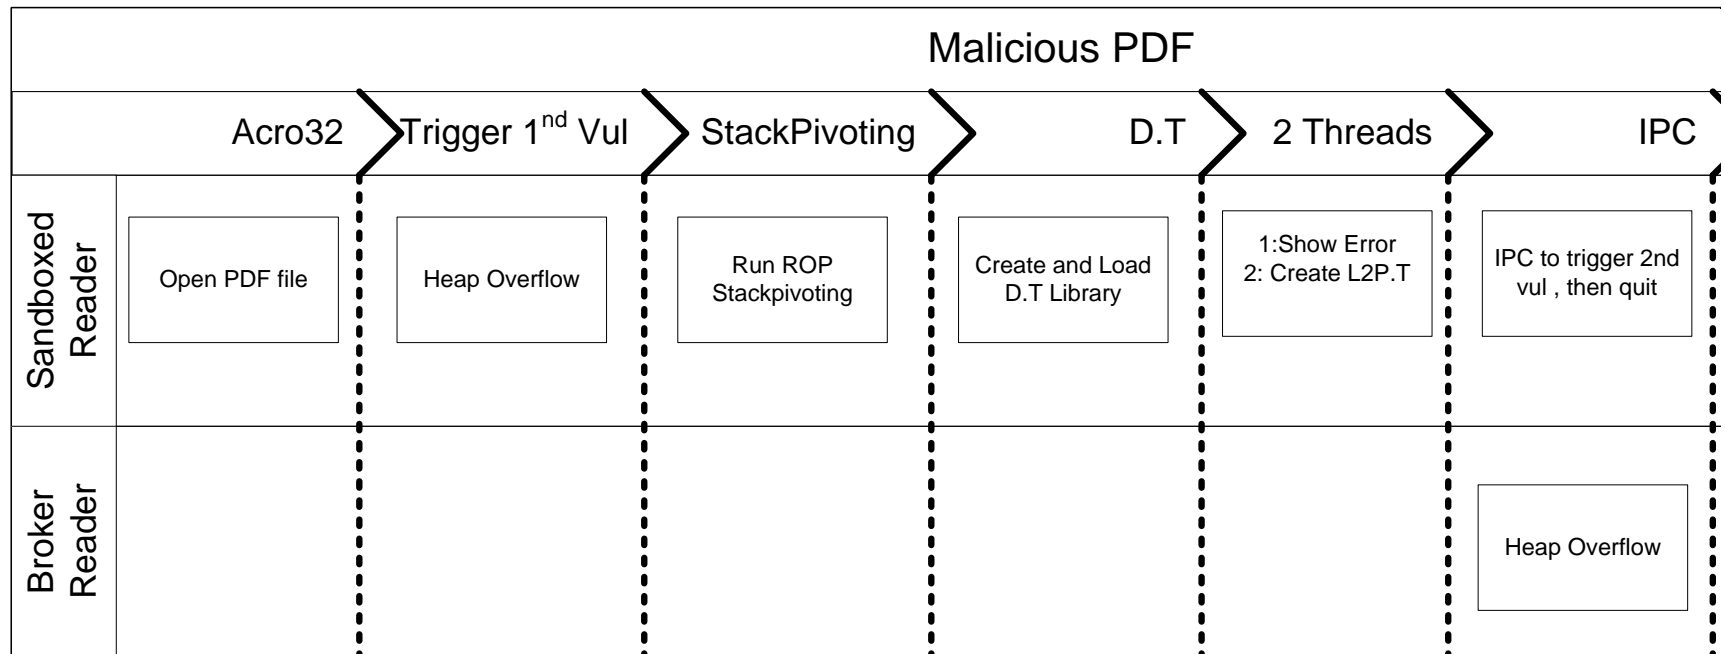
---

# APSA13-02 PDF 0-day

- Reported by FireEye in February 2013
- Best Client-Side Bug - CVE-2013-0641
- sophisticated ROP only without shellcode
- First public in the wild exploit Adobe Sandbox Bypassing



# CVE-2013-0640 Exploit



# CVE-2013-0640 PDF 0-day analysis (1)

## ➤ Trigger Point in AcroForm.api

```
.text:208A54D3  
.text:208A54D3 exploittrigger:  
.text:208A54D3      call    dword ptr [eax]
```

## ➤ Stack Pivoting

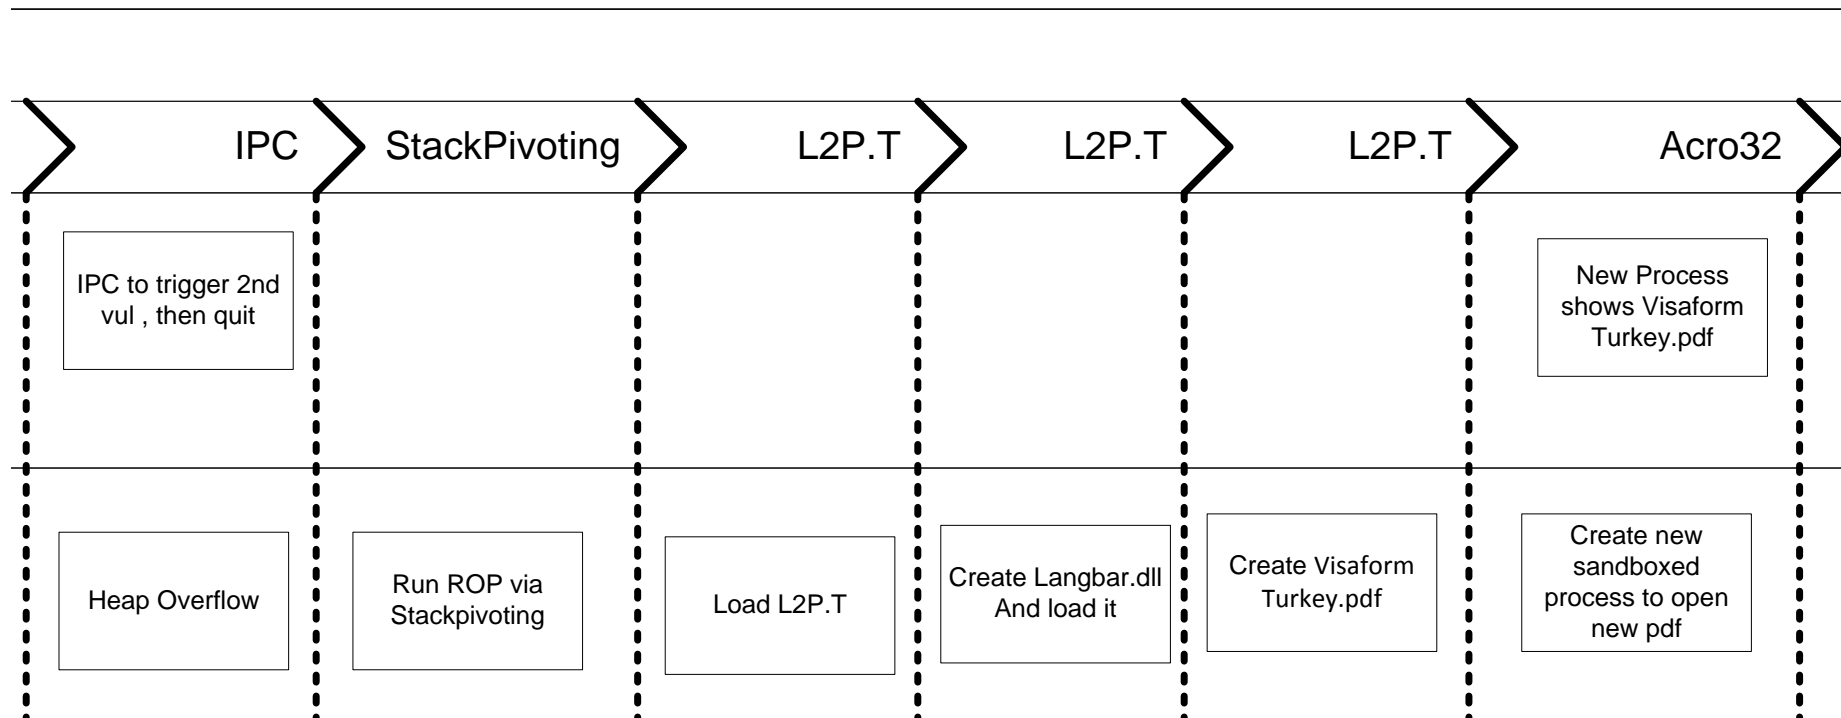
```
text:209B9F42      mov     eax, [ecx+4]  
text:209B9F45      test    eax, eax  
text:209B9F47      jz      short loc_209B9F57  
text:209B9F49      push    eax  
text:209B9F4A      mov     eax, dword_2128C66C  
text:209B9F4F      call    dword ptr [eax+5Ch]  
text:209B9F52      pop     ecx  
text:209B9F53      movzx   eax, ax  
text:209B9F56      retn  
  
.text:209B9F42      mov     eax, [ecx+4]  
.text:209B9F45      test    eax, eax  
.text:209B9F47      jz      short loc_209B9F57  
.text:209B9F49      push    eax  
.text:209B9F4A      mov     eax, dword_2128C66C  
.text:209B9F4A ; -----  
.text:209B9F4F      db      0FFh  
.text:209B9F50 ; -----  
.text:209B9F50      push    eax  
.text:209B9F51      pop     esp  
.text:209B9F52      pop     ecx  
.text:209B9F53      movzx   eax, ax  
.text:209B9F56      retn
```

# CVE-2013-0640 PDF 0-day analysis (2)

---

- Lots of ROP gadgets
  - Get ntdll.dll and get related API address
    - for example, RtlDecompressBuffer and CryptStringToBinaryA
  - Call CryptStringToBinaryA to convert one string to binary
  - Call RtlDecompressBuffer to decompress binary to real D.T binary code in memory
  - Finally, the ROP gadget will call GetTempPathA to get current temp path, it's sandboxed path, create D.T under this path and call LoadLibraryA to run D.T.
- D.T dll will trigger 2nd 0-day to load L2P.T in broker process

# CVE-2013-0641 Exploit





# CVE-2013-0641 PDF 0-day analysis (1)

---

## ➤ Trigger Point in acro32.exe

```
.text:0049728A secondvul_triggerpoint:  
.text:0049728A          call     eax
```

## ➤ Stack Pivoting

```
.text:6F881564 ; -----  
.text:6F881565          db  0FFh  
.text:6F881566 ; -----  
.text:6F881566          push    edx  
.text:6F881567          pop     esp  
.text:6F881568          pop     ebp  
.text:6F881569          retn    8
```

## ➤ ROP Gadget will load L2P.T and kick off malware infection

---

# Demo

---

# Misc

---

- Interrupted EIP
    - Sometimes the Interrupted EIP is different from target EIP
  - Evasion
    - Dummy exploit to trigger code to avoid detection before real exploit code happen
  - Global Events
    - PMU Events are global
  - Missed Event
    - PMU is designed for profiling, not for security
  - LBR Record
  - User Mode Scheduler
-

# Interrupt EIP and Event EIP

- The Interrupt could happen on different EIP from Event EIP

FromAddr=0x5792d826

IntAddr=0x57932ba3

ToAddr=0x57932b80

5792d807 8b4c2428	mov	ecx,dword ptr [esp+28h]	57932b80 83ec08	sub	esp,8
5792d80b e840faffff	call	Flash32_15_0_0_189+0xbd250 (5792d250)	57932b83 53	push	ebx
5792d810 8bd0	mov	edx,eax	57932b84 8b5c2414	mov	ebx,dword ptr [esp+14h]
5792d812 2bd6	sub	edx,esi	57932b88 55	push	ebp
5792d814 8d3c97	lea	edi,[edi+edx*4]	57932b89 8b6c2414	mov	ebp,dword ptr [esp+14h]
5792d817 8bf0	mov	esi,eax	57932b8d 807d2c01	cmp	byte ptr [ebp+2Ch],1
5792d819 eb14	jmp	Flash32_15_0_0_189+0xbd82f (5792d82f)	57932b91 56	push	esi
5792d81b 8d4c2420	lea	ecx,[esp+20h]	57932b92 8b742420	mov	esi,dword ptr [esp+20h]
5792d81f 51	push	ecx	57932b96 57	push	edi
5792d820 53	push	ebx	57932b97 0f85ae000000	jne	Flash32_15_0_0_189+0xc2c4b (57932c4b)
5792d821 52	push	edx	57932b9d 8b4d14	mov	ecx,dword ptr [ebp+14h]
5792d822 8b5038	mov	edx,dword ptr [eax+38h]	57932ba0 8b7d50	mov	edi,dword ptr [ebp+50h]
5792d825 50	push	eax	57932ba3 8bc3	mov	eax,ebx
5792d826 ffd2	call	edx	57932ba5 c1e010	shl	eax,10h
5792d828 83c410	add	esp,10h	57932ba8 89442410	mov	dword ptr [esp+10h],eax
5792d82b 83ef04	sub	edi,4	57932bac 8b5140	mov	edx,dword ptr [ecx+40h]
5792d82e 4e	dec	esi	57932baf 8d442410	lea	eax,[esp+10h]
5792d82f 8b5508	mov	edx,dword ptr [ebp+8]	57932bb3 c1e210	shl	edx,10h
5792d832 85f6	test	esi,esi	57932bb6 50	push	eax
5792d834 7fba	jb	Flash32_15_0_0_189+0xbd7f0 (5792d7f0)	57932bb7 8bc8	mov	ecx,eax
			57932bb9 89542418	mov	dword ptr [esp+18h],edx
			57932bbd 51	push	ecx
			57932bbe 8d5724	lea	edx,[edi+24h]
			57932b9d 8b4d14	mov	ecx,dword ptr [ebp+14h]
			57932ba0 8b7d50	mov	edi,dword ptr [ebp+50h]
			57932ba3 8bc3	mov	eax,ebx
			57932ba5 c1e010	shl	eax,10h
			57932ba8 89442410	mov	dword ptr [esp+10h],eax

# Repeated Indirect Call Misprediction

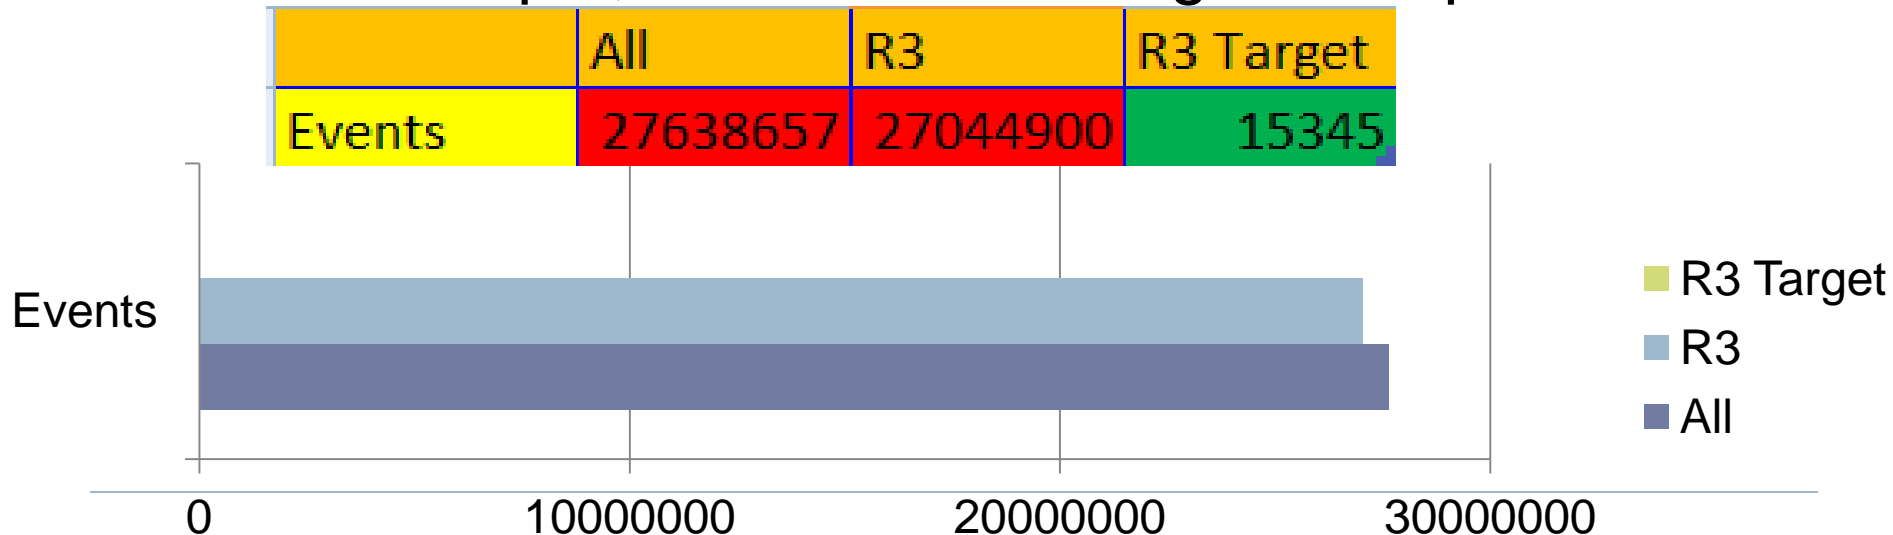
---

4:41:40.238	PM	FromAddr=0x7612c498 , IntAddr=0x759a68ef , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c171 , IntAddr=0x762bc470 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7610cc50 , IntAddr=0x759a68ef , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c847 , IntAddr=0x759a68f2 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612d0b3 , IntAddr=0x762bc470 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761bc426 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761bc453 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7610d9e7 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7610e0a1 , IntAddr=0x7610e0a8 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761465f3 , IntAddr=0x759a68f2 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7610dc15 , IntAddr=0x759a68f2 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c7d4 , IntAddr=0x7612c7db , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c1c2 , IntAddr=0x759a68f2 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c1f8 , IntAddr=0x759a68f2 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c192 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7610db94 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c22b , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c259 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x76115672 , IntAddr=0x762bc478 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c043 , IntAddr=0x762bc470 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7610d68b , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7610d248 , IntAddr=0x77942dd6 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761bc9f2 , IntAddr=0x77942dd6 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761bc426 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7610d9af , IntAddr=0x77942dd6 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761bc854 , IntAddr=0x77942dd6 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c7d4 , IntAddr=0x762bc478 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c1c2 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x76115120 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761150c6 , IntAddr=0x762bc472 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761bc426 , IntAddr=0x762bc478 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c1c2 , IntAddr=0x762bc470 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c1f8 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7612c192 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x7610db94 , IntAddr=0x759a68f5 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761bc9f2 , IntAddr=0x77942de0 , ToAddr=0x762bfd8d
4:41:40.238	PM	FromAddr=0x761bc854 , IntAddr=0x77942dd6 , ToAddr=0x762bfd8d

---

# Global Events

- Without MSR enabling/disabling during OS Scheduler
  - The event could happen on every code on logical processor whatever the process
- With proper MSR enabling/disabling during context switch, performance will be improved
  - For example, 5-6 events on target CR3 per 10K



# LBR Record

- Must freeze LBRs on PMI
- MSR
  - MSR\_LASTBRANCH\_TOS
  - MSR\_LASTBRANCH\_X\_FROM\_IP
  - MSR\_LASTBRANCH\_X\_TO\_IP

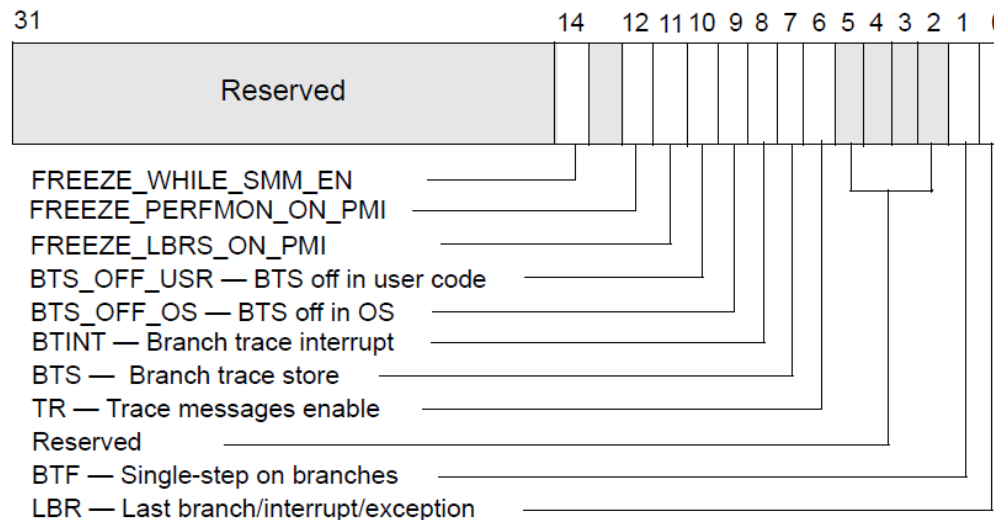


Figure 17-3. IA32\_DEBUGCTL MSR for Processors based on Intel Core microarchitecture

# MSR\_LASTBRANCH\_X\_FROM\_IP/MSR\_LASTBRANCH\_X\_TO\_IP

---

- Silvermont

- MSR\_LASTBRANCH\_0\_FROM\_IP (address 40H)
- MSR\_LASTBRANCH\_0\_TO\_IP (address 60H)

- Haswell

- MSR\_LASTBRANCH\_0\_FROM\_IP (address 680H)
  - MSR\_LASTBRANCH\_0\_TO\_IP (address 6C0H)
-



# Bonus

---

- Improved LIFO filter to capture Call Stack
  - Block simple repeated LBR FLUSH

**Table 17-11. MSR\_LBR\_SELECT for Intel® microarchitecture code name Haswell**

Bit Field	Bit Offset	Access	Description
CPL_EQ_0	0	R/W	When set, do not capture branches occurring in ring 0
CPL_NEQ_0	1	R/W	When set, do not capture branches occurring in ring >0
JCC	2	R/W	When set, do not capture conditional branches
NEAR_REL_CALL	3	R/W	When set, do not capture near relative calls
NEAR_IND_CALL	4	R/W	When set, do not capture near indirect calls
NEAR_RET	5	R/W	When set, do not capture near returns
NEAR_IND_JMP	6	R/W	When set, do not capture near indirect jumps except near indirect calls and near returns
NEAR_REL_JMP	7	R/W	When set, do not capture near relative jumps except near relative calls.
FAR_BRANCH	8	R/W	When set, do not capture far branches
EN_CALLSTACK <sup>1</sup>	9		Enable LBR stack to use LIFO filtering to capture Call stack profile
Reserved	63:10		Must be zero

# Summary

---

- Most Exploits will cause branch misprediction with unintended/intended code
  - Branch mispredicted events are useful to detect exploits with minimized performance impacts
  - APSA13-2 exploit was successfully detected by branch mispredicted based approach
-

# Thanks!

---



Xiaoning.li@intel.com  
mcrouse@seas.harvard.edu

*Acknowledge Haifei Li and Dave Marcus' review!*

---

# Reference

---

- [1] White Phosphorus Exploit Pack Sayonara ASLR DEP Bypass Technique, <http://www.whitephosphorus.org/sayonara.txt>
  - [2] Active Zero-Day Exploit Targets Internet Explorer Flaw, <http://blogs.mcafee.com/mcafee-labs/active-zero-day-exploit-targets-internet-explorer-flaw>
  - [3] Smashing the Heap with Vector: Advanced Exploitation Technique in Recent Flash Zero-day Attack, Haifei Li
  - [4] Interpreter Exploitation: Pointer Inference and JIT Spraying, Dion Blazakis, BlackHat DC 2010
  - [5] Attacking the Windows 7/8 Address Space Randomization, <http://kingcope.wordpress.com/2013/01/24/attacking-the-windows-78-address-space-randomization>
  - [6] DEP/ASLR bypass without ROP/JIT, Yang Yu, CanSecWest 2013
  - [7] The BlueHat Prize, <http://www.microsoft.com/security/bluehatprize>
  - [8] kBouncer: Efficient and Transparent ROP Mitigation, Vasilis Pappas
  - [9] Taming the ROPE on Sandy Bridge, Georg Wicherski, SysCan 2013
  - [10] Security Breaches as PMU Deviation: Detecting and Identifying Security Attacks Using Performance Counters, Liwei Yuan etc, APSYS 2011
  - [11] Intel® 64 and IA-32 Architectures Software Developer's Manual Combined Volumes:1, 2A, 2B, 2C, 3A, 3B, and 3C, [www.intel.com](http://www.intel.com)
  - [12] PRACTICAL RETURN-ORIENTED PROGRAMMING, Dino Dai Zovi, RSA 2010
  - [13] The Past, Present, and Future of Flash Security, Haifei Li, SysCan360 2012
  - [14] return oriented exploitation, Dino Dai Zovi, Blackhat 2010
  - [15] Security Mitigations for Return-Oriented Programming Attacks, Piotr Bania, [http://piotrbania.com/all/articles/pbania\\_rop\\_mitigations2010.pdf](http://piotrbania.com/all/articles/pbania_rop_mitigations2010.pdf)
-

# Backup

---

---

# Architectural Performance Monitoring Version 1

---

- Bit field layout of IA32\_PERFECTSELx is consistent across micro architectures
  - Addresses of IA32\_PERFECTSELx MSRs remain the same across micro architectures
  - Addresses of IA32\_PMC MSRs remain the same across micro architectures
  - Each logical processor has its own set of IA32\_PERFECTSELx and IA32\_PMCx MSRs.  
Configuration facilities and counters are not shared between logical processors sharing a processor core
-

# Architectural Performance Monitoring Version 1 Facilities

---

- IA32\_PMCx MSRs start at address 0C1H and occupy a contiguous block of MSR address space the number of MSRs per logical processor is reported using CPUID.0AH:EAX[15:8]
  - IA32\_PERFECTSELx MSRs start at address 186H and occupy a contiguous block of MSR address space.
  - The bit width of an IA32\_PMCx MSR is reported using the CPUID.0AH:EAX[23:16]. This the number of valid bits for read operation. On write operations, the lower-order 32 bits of the MSR may be written with any value, and the high-order bits are sign-extended from the value of bit 31
  - Bit field layout of IA32\_PERFECTSELx MSRs is defined architecturally
-